ESnet

ENERGY SCIENCES NETWORK

Modern Cyberinfrastructure: The Ladder to the Shoulders Of Giants

Eli Dart, Science Engagement

Energy Sciences Network (ESnet)

Lawrence Berkeley National Laboratory

PEARC20

Virtual (Coronapocalypse)

July 27, 2020

U.S. DEPARTMENT OF ENERGY

Office of Science

BERKELEY LAB

Our Community Has Made Great Progress

- Over the past 8-10 years or so, we have made great strides forward
- Science networks are big, fast, and clean
  - High speed regional and national networks
  - R&E exchanges
  - Campus networks
  - International connectivity
- Science networks are instrumented for performance
  - perfSONAR
  - Critical for ensuring correct operation
  - Invaluable for timely resolution of problems

ESnet

Our Community Has Made Great Progress

- The Science DMZ model is widely deployed
  - Campuses, laboratories, experiments
  - HPC facilities
  - Some data portals (more on this later)
- DTNs in the Science DMZs
  - Connect storage to high speed networks
  - HPC filesystems
  - Experiment data acquisition systems
- Data orchestration platforms running on the DTNs
  - This is what the scientist sees
  - Capable platforms allow orchestration rather than clunky user-driven scripting or manual downloads

ESnet

Data Ecosystem - Abstract Network Diagram

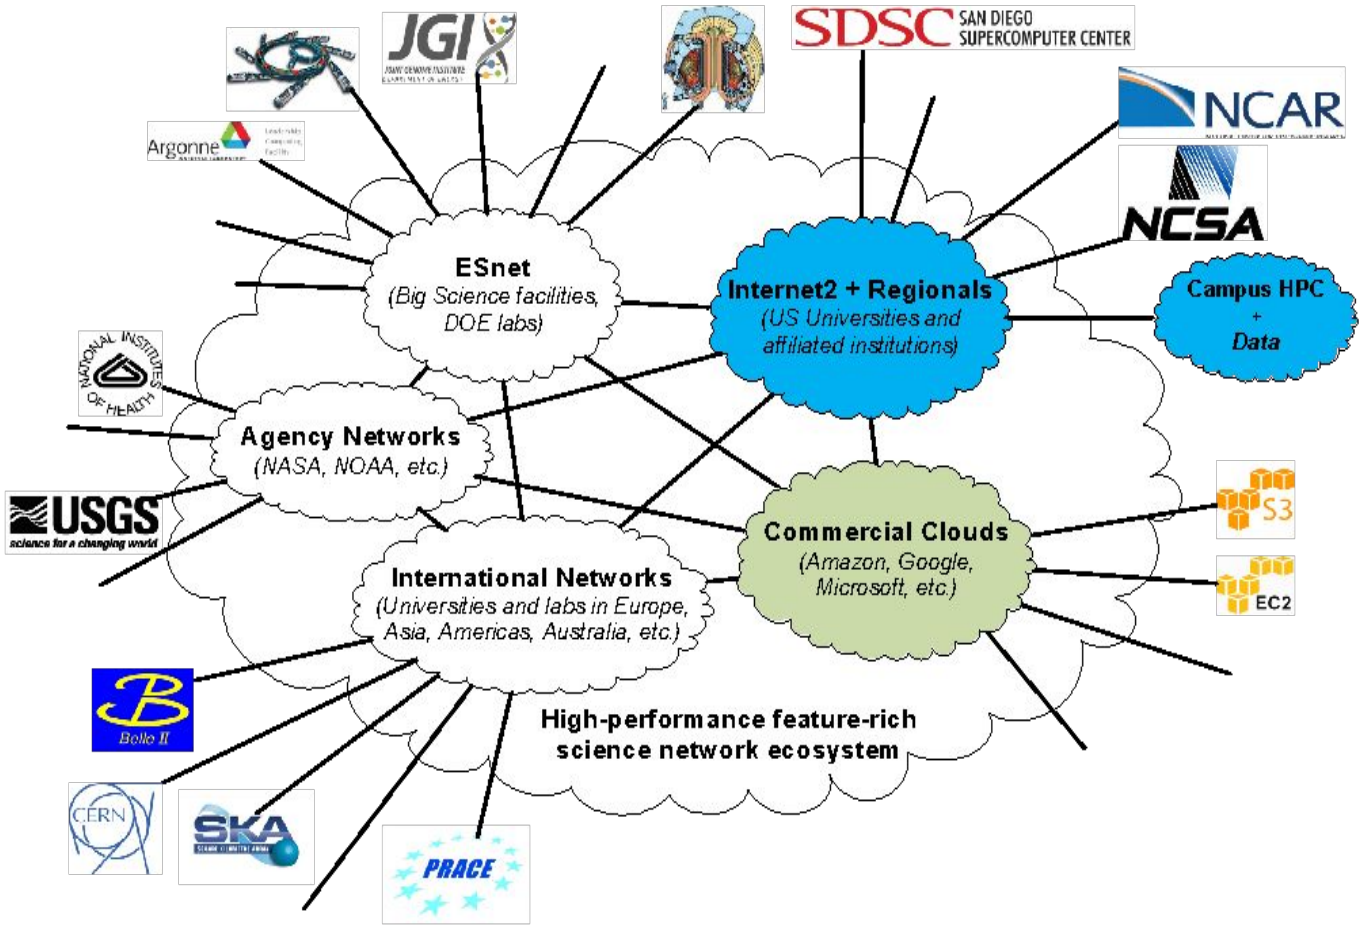

ESnet

Data Ecosystem – Block Visualization

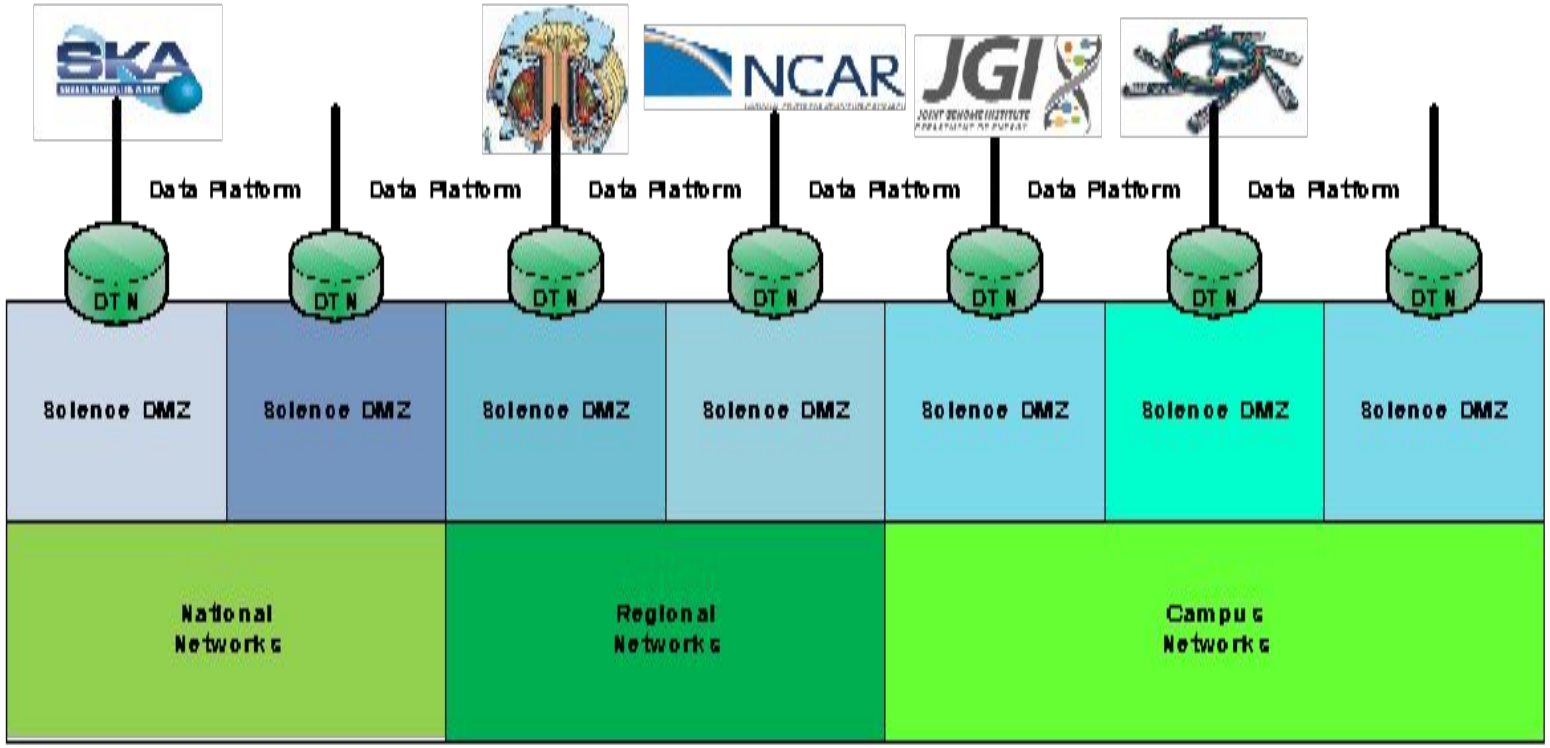

ESnet

Data Ecosystem – Concentric View

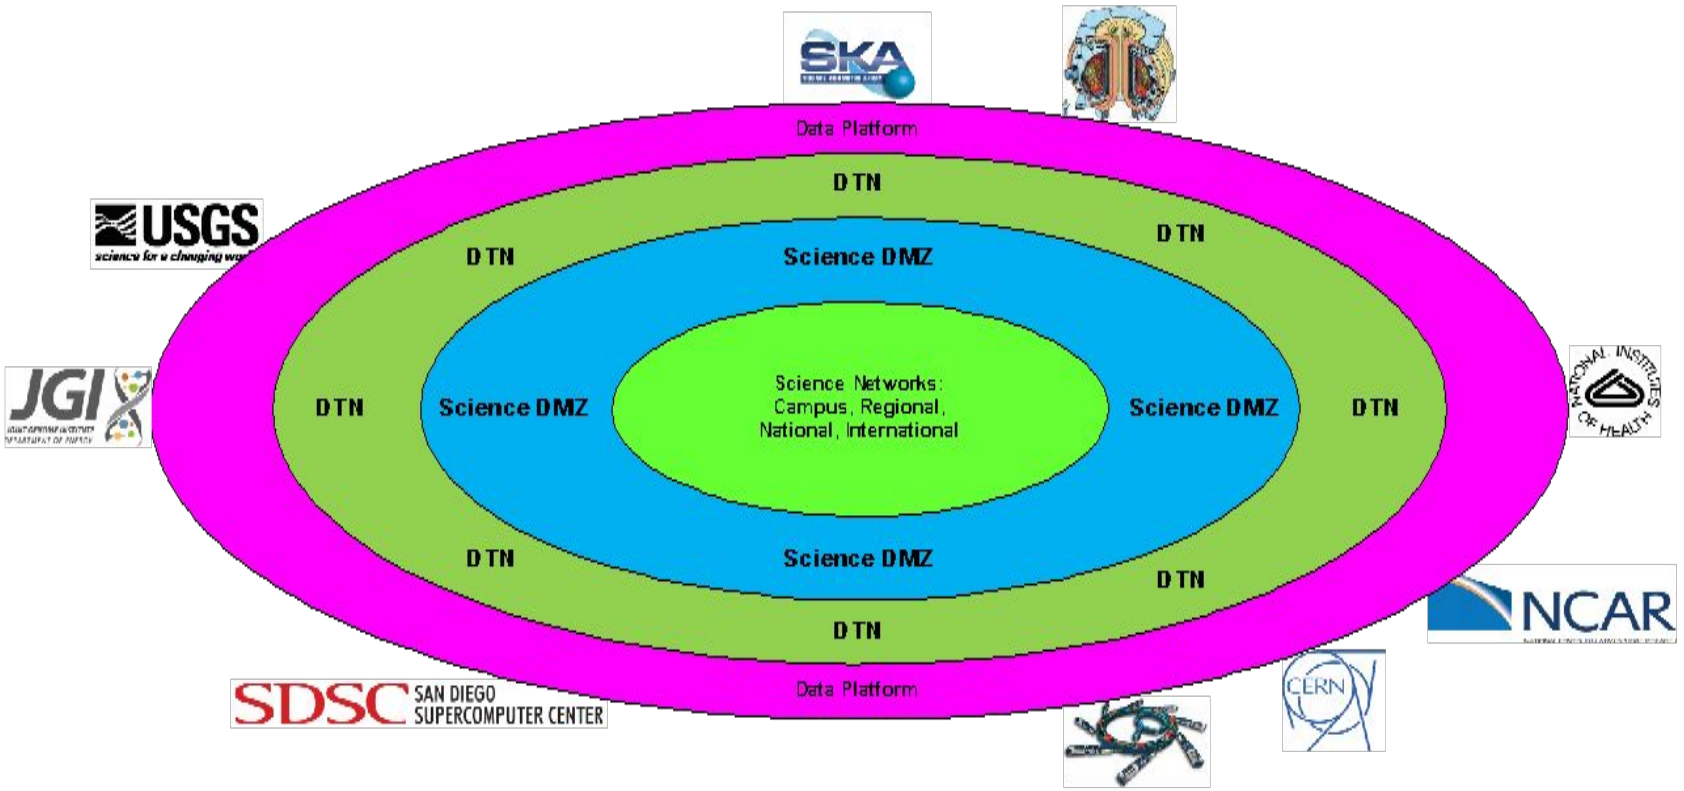

ESnet

What Remains To Be Done?

- We aren't all the way there yet (unfortunately)
- The diagrams show a vision that is not yet fully realized
- Three major tasks remain
  - Deployment of an interoperable platform across Science DMZs
    - This includes test, verification, and performance engineering
    - Partially complete
  - Integrating the major data repositories and portals with the platform
    - This has begun – lots left to do
  - Onboarding scientists and collaborations
    - Science Engagement
    - We understand it, but we need to scale it
- Remember – this has to be useful to scientists, so it has to work for them

ESnet

Interoperable Platform Deployment

- This is partially complete
- Necessary features
  - Automation
  - Fault recovery
  - Data integrity
  - Integration with web-based portals
- Several platforms exist
  - Globus
    - Significant deployment in NSF and DOE spaces
    - Basis for examples shown here
  - XRootD (LHC experiments)
  - OSG Stack
- Key point – **the scientist must not be made the integrator**
  - If the scientist is the integrator, they will use HTTP and rsync+SSH forever
  - The old tools don't scale, but the scientists can't build the better platforms themselves
  - **WE** MUST DO THIS

ESnet

Example Of Platform Power (Petascale DTN)

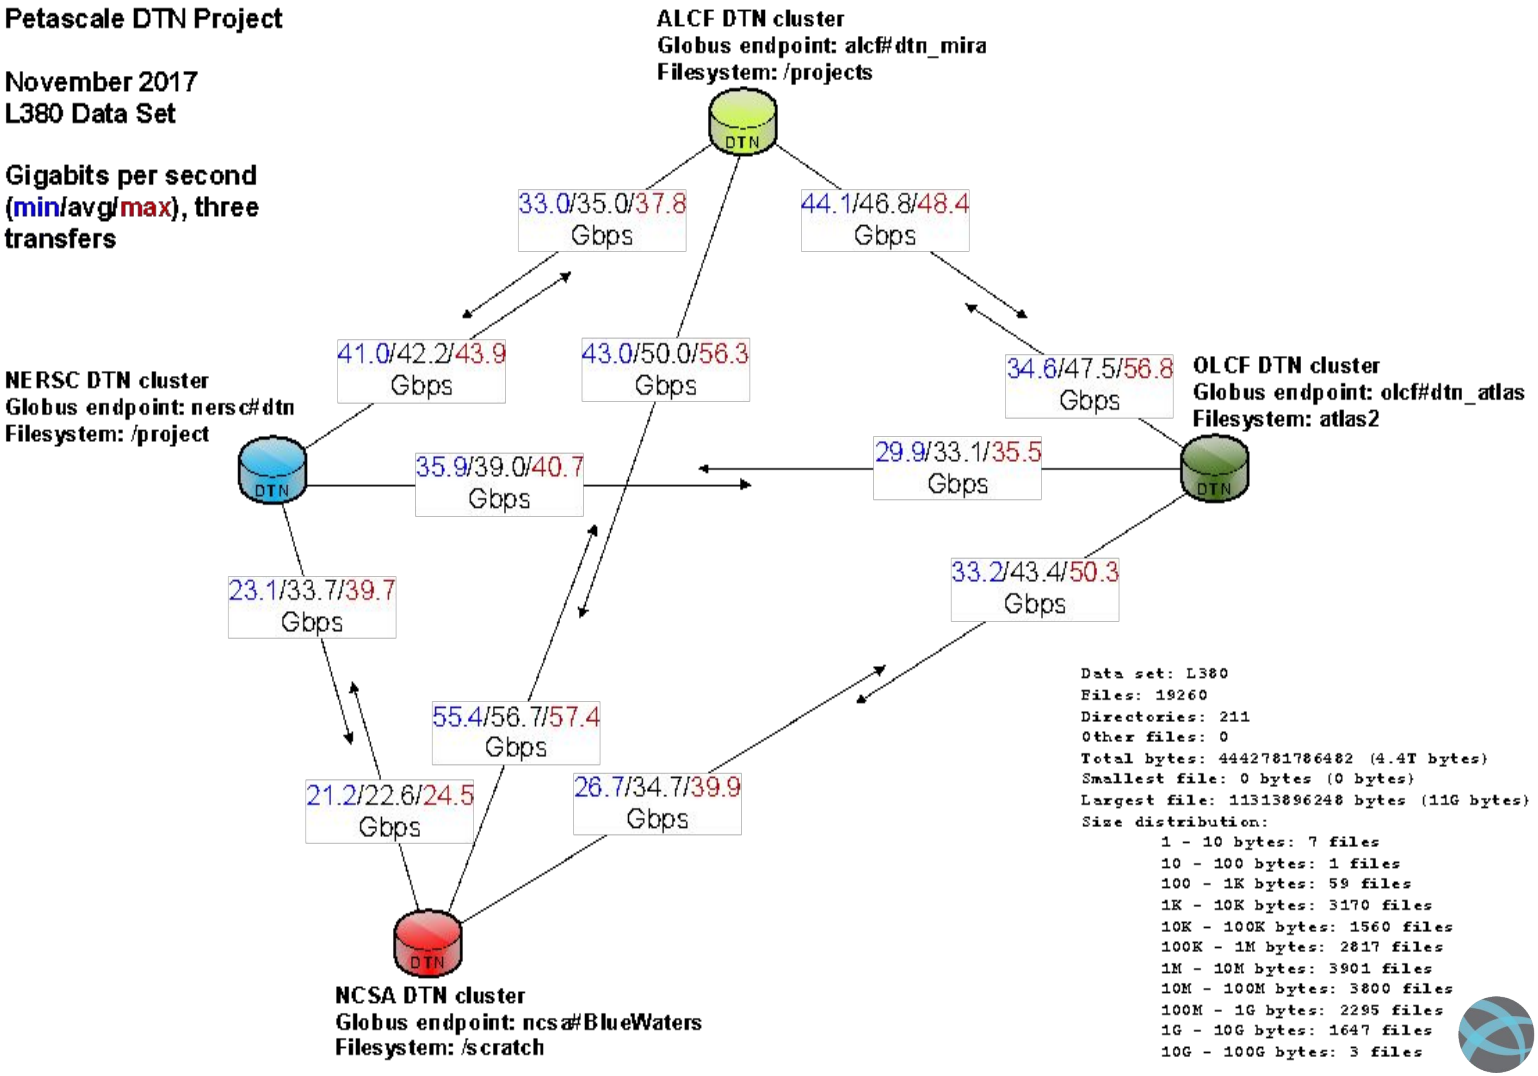

Science DMZ – HPC Center DTN Cluster

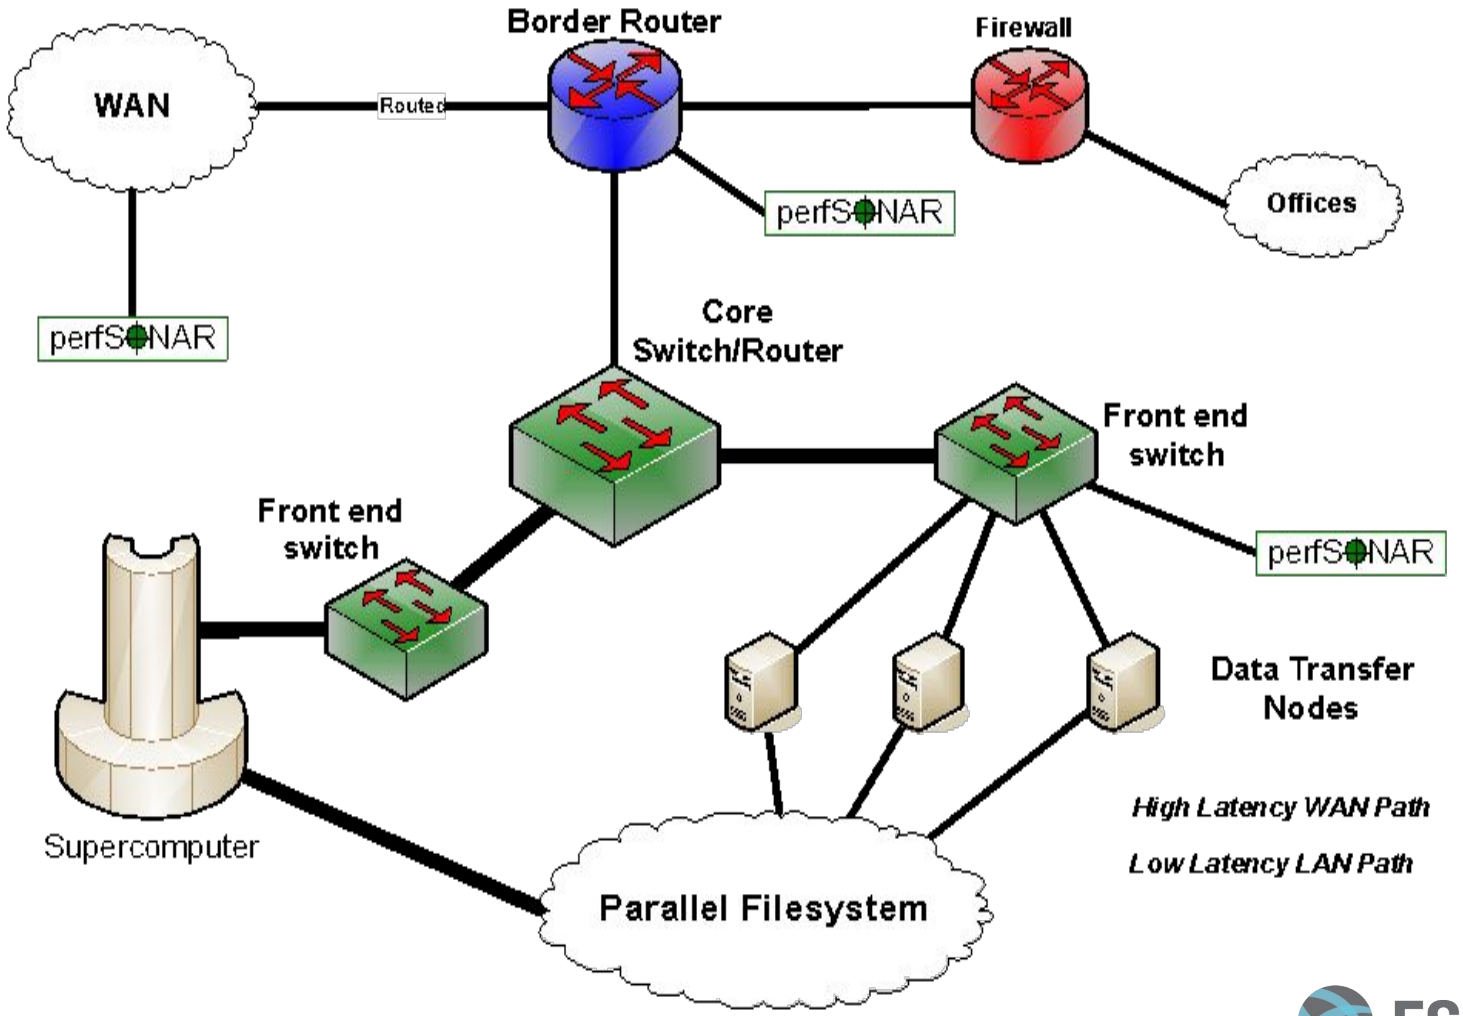

Science Data Portals

- Large repositories of scientific data
  - Climate data
  - Sky surveys (astronomy, cosmology)
  - Many others
  - Data search, browsing, access
- Many scientific data portals were designed 15+ years ago
  - Single-web-server design
  - Data browse/search, data access, user awareness all in a single system
  - All the data goes through the portal server
    - In many cases by design
    - E.g. embargo before publication (enforce access control)
  - Better than old command-line FTP, but outdated by today's standards

ESnet

Legacy Portal Design

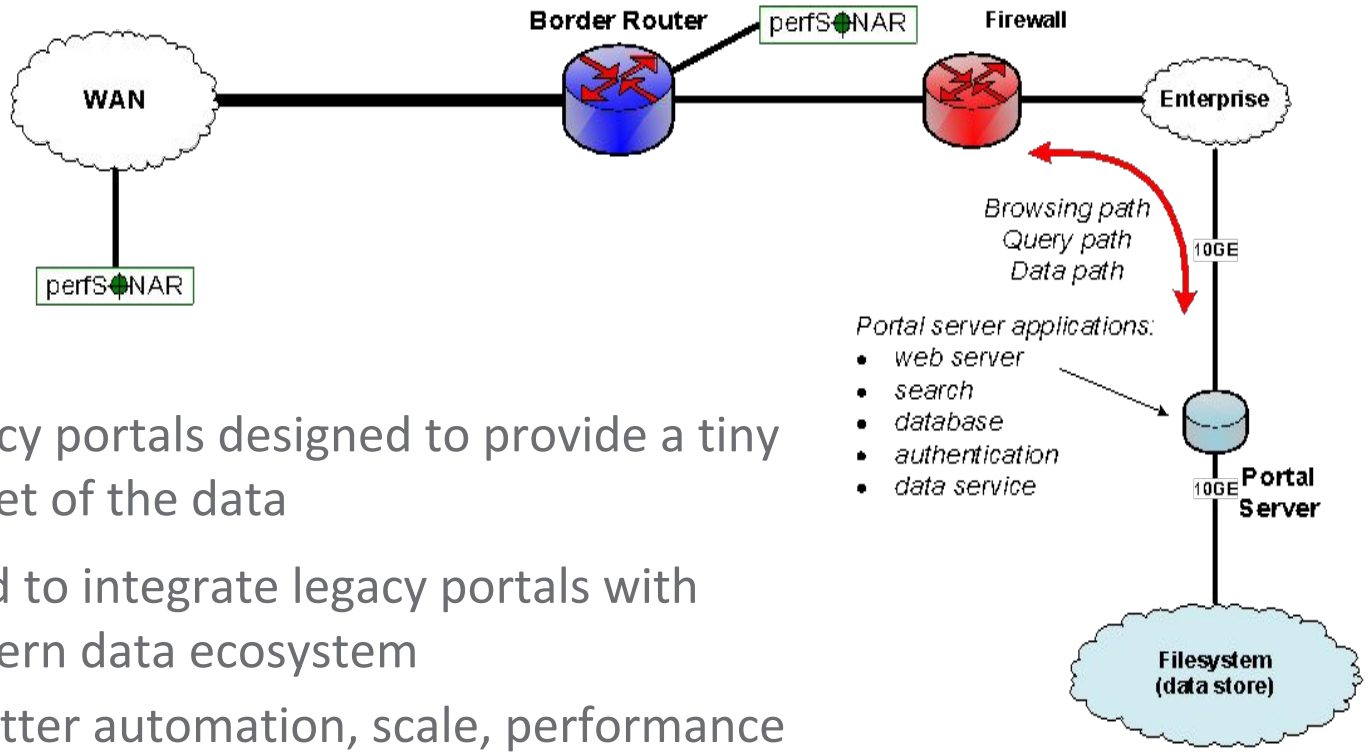

- Legacy portals designed to provide a tiny subset of the data
- Need to integrate legacy portals with modern data ecosystem
  - Better automation, scale, performance
  - Connectivity to HPC

ESnet

Next-Generation Portal Leverages Science DMZ

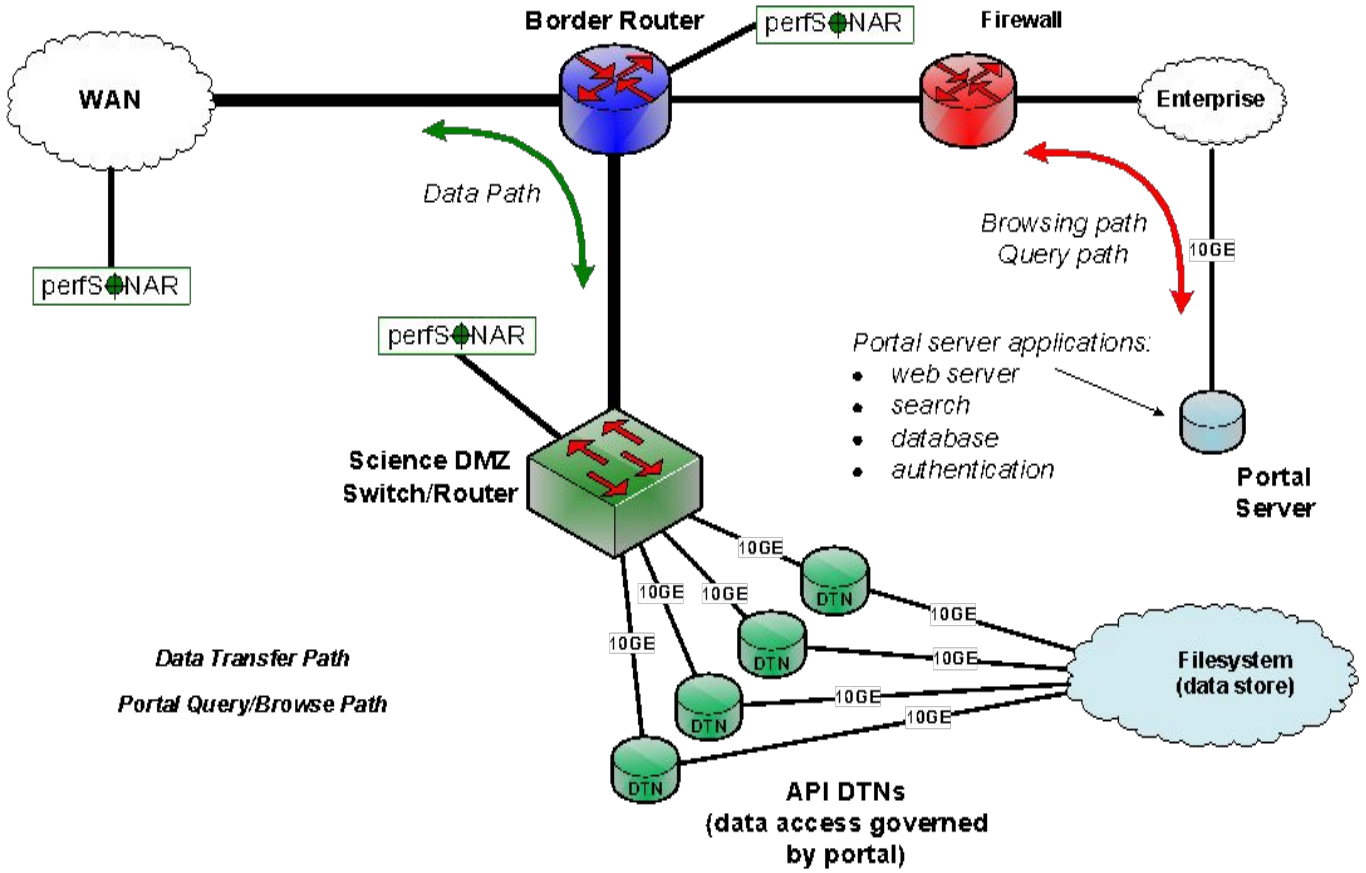

https://peerj.com/articles/cs-144/

ESnet

JGI Data Portal

Searching for Projects

- Explore what you can do here.
- Search projects/proposals using “Advanced Search” filters.

Downloading Files

- Download over the web
- Download large number of files with Globus service.
- Download via API using scripting or programming
- Download with "Cart" by collecting projects/portals of your interest.

Looking for Access

- Looking for data and do not have access to the private portal? Please contact PI
- How to grant access to your proposal/project/genome? Get Instructions.

JGI Genome Portal

New Feature: "Bulk Downloads"

collect your favorite projects and download them in **bulk** with our new feature **"CART"**. Ability to download files with Portal or via Globus.

Find out more details

What's New

New feature: "Download with Cart" 🛒

A convenient way to collect projects/genomes/metagenomes of your interest and download all files associated with them in **bulk**.
Read more and provide your comments and suggestions for this feature to our team.

My Favorites

My Favorites: New Feature - Based on Your Feedback

This feature allows to save your filtered search results to "My Favorites" and access it later.

The "Tree of Life"

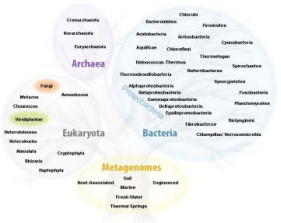

Please use our powerful search or go to the "Tree of Life" if it is the most convenient way for you to reach your genomes/projects.

ESnet

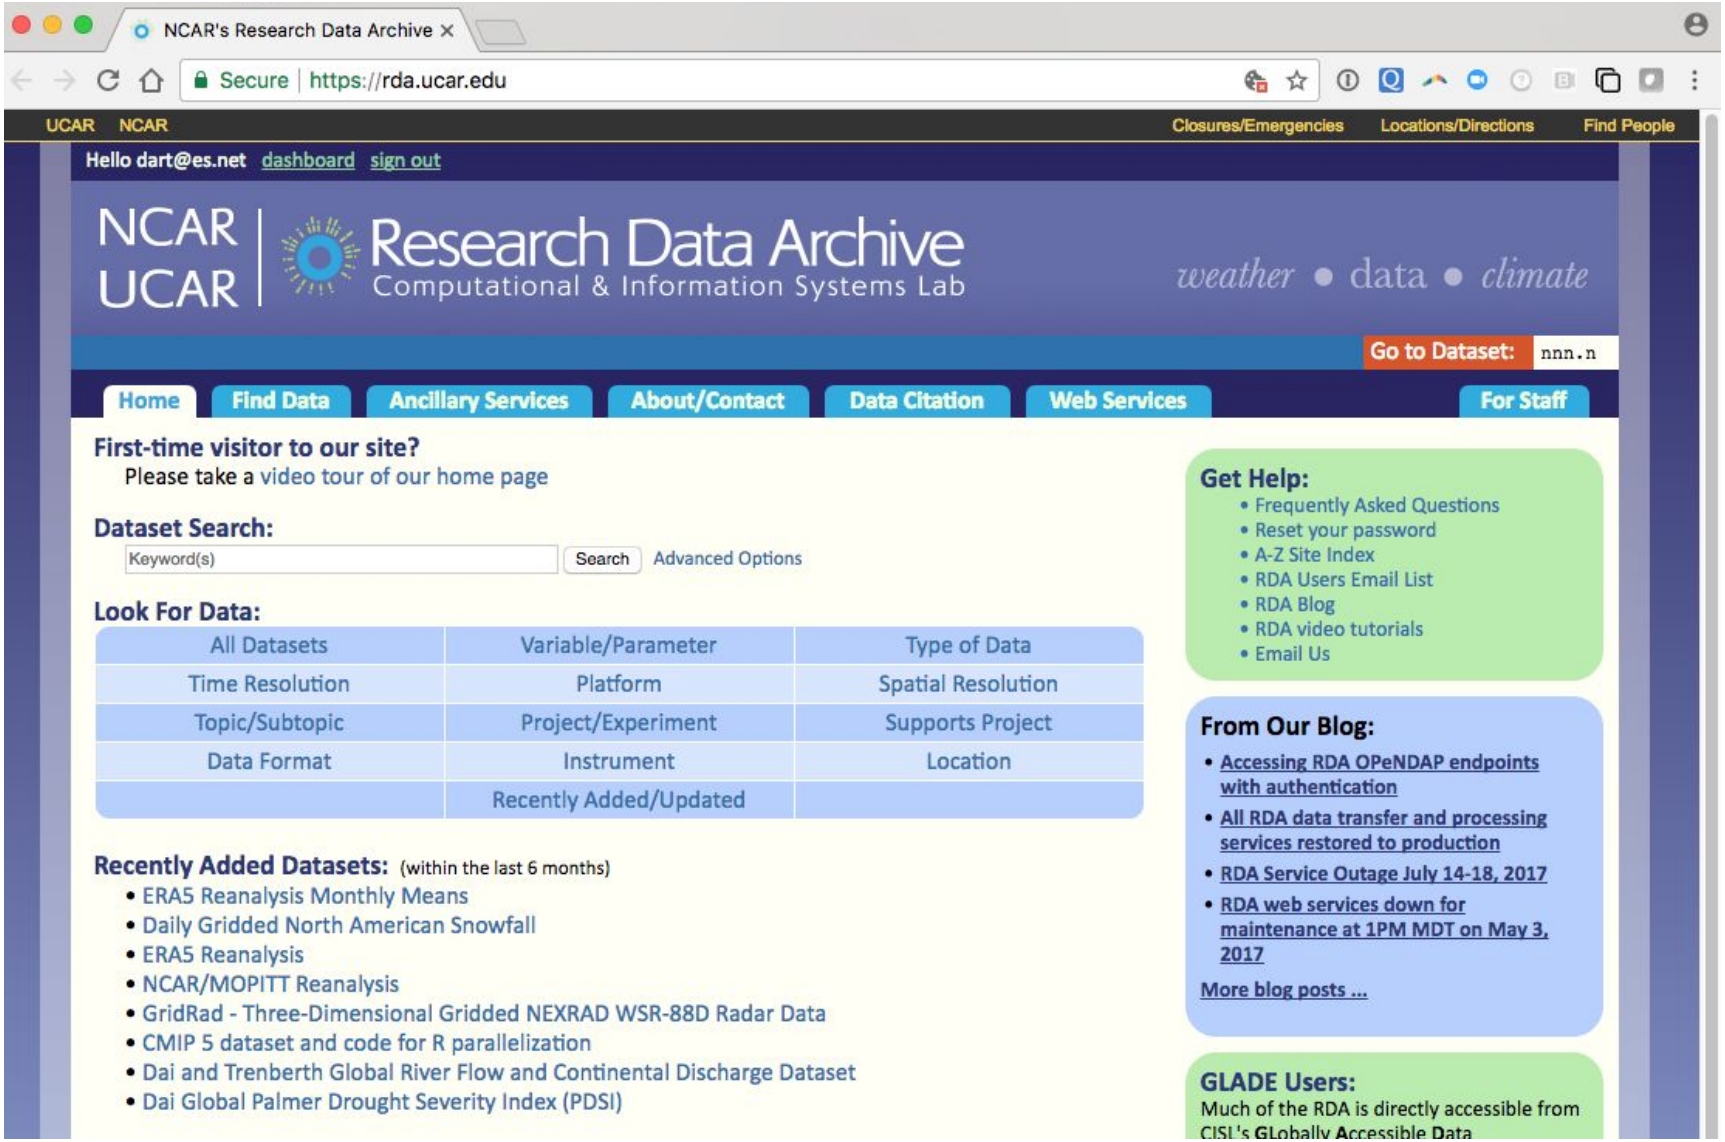
NCAR's Research Data Archive
Secure | https://rda.ucar.edu
UCAR NCAR
Closures/Emergencies Locations/Directions Find People
Hello dart@es.net dashboard sign out
NCAR UCAR
Research Data Archive
Computational & Information Systems Lab
weather • data • climate
Go to Dataset: nnn.n
Home Find Data Ancillary Services About/Contact Data Citation Web Services For Staff
First-time visitor to our site?
Please take a video tour of our home page
Dataset Search:
Keyword(s) Search Advanced Options
Look For Data:
All Datasets Variable/Parameter Type of Data
Time Resolution Platform Spatial Resolution
Topic/Subtopic Project/Experiment Supports Project
Data Format Instrument Location
Recently Added/Updated
Recently Added Datasets: (within the last 6 months)
• ERA5 Reanalysis Monthly Means
• Daily Gridded North American Snowfall
• ERA5 Reanalysis
• NCAR/MOPITT Reanalysis
• GridRad - Three-Dimensional Gridded NEXRAD WSR-88D Radar Data
• CMIP 5 dataset and code for R parallelization
• Dai and Trenberth Global River Flow and Continental Discharge Dataset
• Dai Global Palmer Drought Severity Index (PDSI)
Get Help:
• Frequently Asked Questions
• Reset your password
• A-Z Site Index
• RDA Users Email List
• RDA Blog
• RDA video tutorials
• Email Us
From Our Blog:
• Accessing RDA OPeNDAP endpoints with authentication
• All RDA data transfer and processing services restored to production
• RDA Service Outage July 14-18, 2017
• RDA web services down for maintenance at 1PM MDT on May 3, 2017
More blog posts ...
GLADE Users:
Much of the RDA is directly accessible from

NCAR RDA Performance to DOE HPC Facilities

- 1.5TB data set
- 1121 files

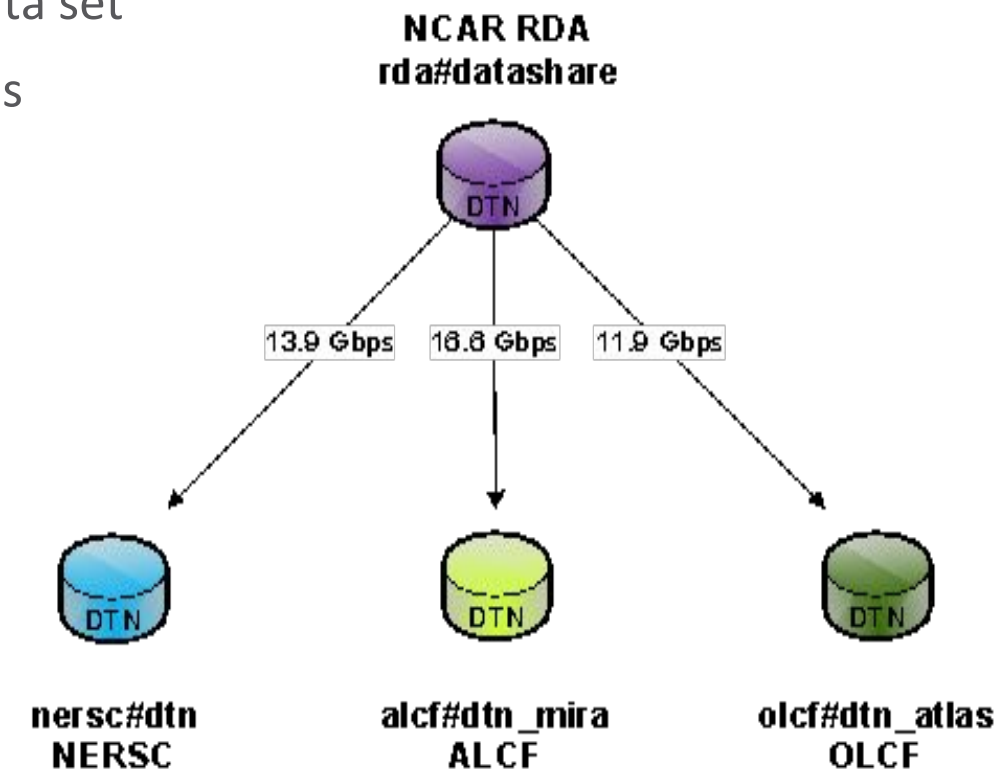

ESnet

Reasons To Scale Data Portals

- Some reasons are obvious
  - Increase in size of data objects (MB □ GB □ 100s of GB)
  - Number of data objects (many thousands per data set)
- Other reasons are paradigm shifts
  - Modern data analysis on HPC can use a *lot* of data
  - Today's HPC facilities are far more capable than in the past
- Retrofit / rebuilt data portals and data repositories
  - Significant wins from increased data analysis

ESnet

Science at Scale: Genomics

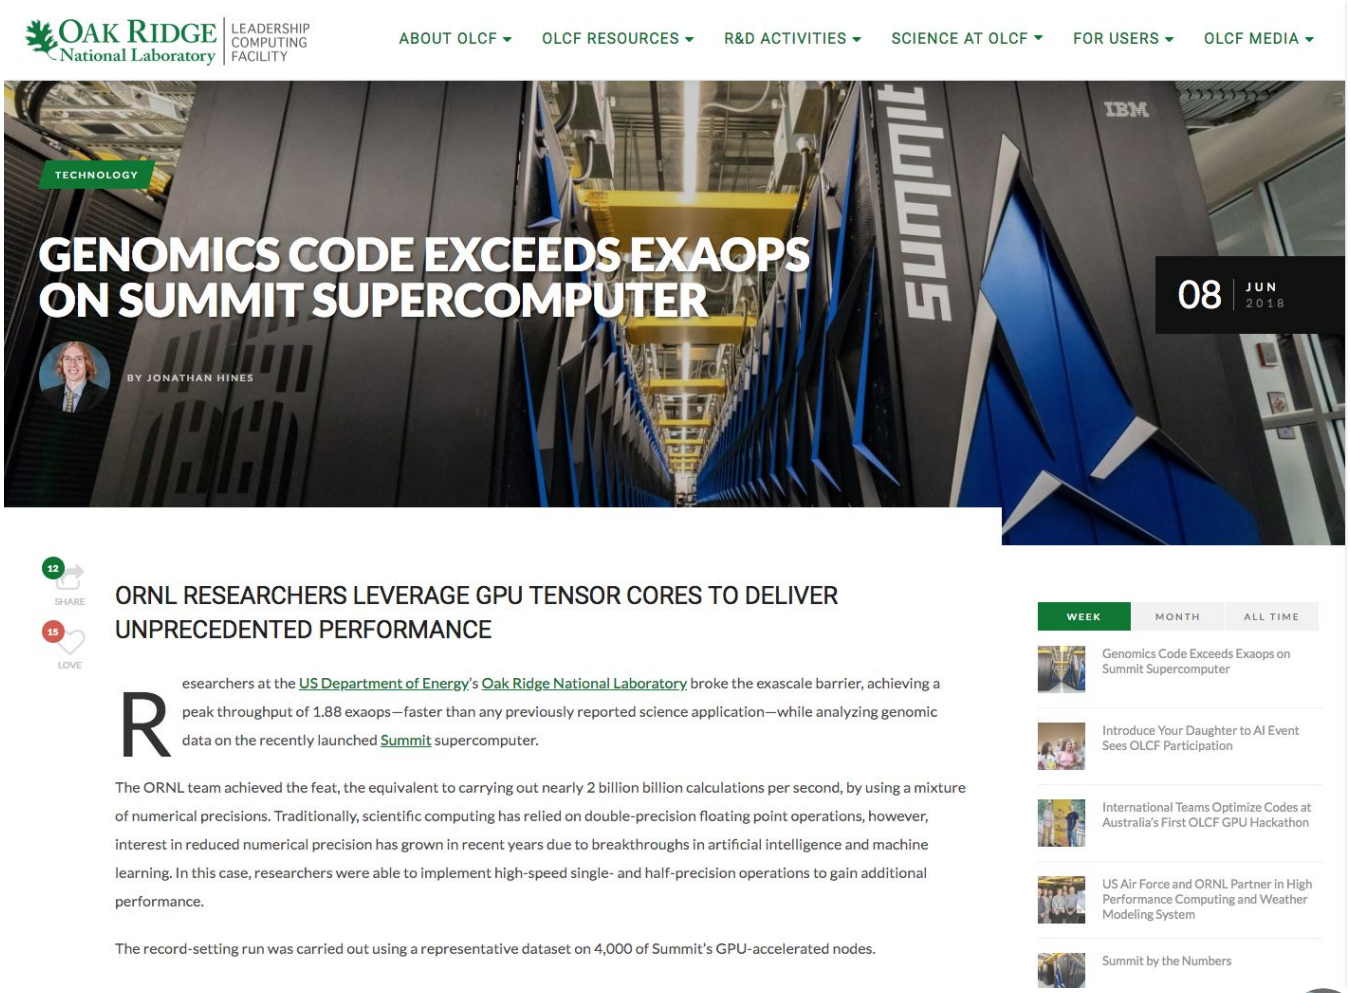

OAK RIDGE National Laboratory | LEADERSHIP COMPUTING FACILITY

ABOUT OLCF ▾ OLCF RESOURCES ▾ R&D ACTIVITIES ▾ SCIENCE AT OLCF ▾ FOR USERS ▾ OLCF MEDIA ▾

TECHNOLOGY

GENOMICS CODE EXCEEDS EXAOPS ON SUMMIT SUPERCOMPUTER

BY JONATHAN HINES

08 JUN 2018

12 SHARE

15 LOVE

ORNL RESEARCHERS LEVERAGE GPU TENSOR CORES TO DELIVER UNPRECEDENTED PERFORMANCE

Researchers at the US Department of Energy's Oak Ridge National Laboratory broke the exascale barrier, achieving a peak throughput of 1.88 exaops—faster than any previously reported science application—while analyzing genomic data on the recently launched Summit supercomputer.

The ORNL team achieved the feat, the equivalent to carrying out nearly 2 billion billion calculations per second, by using a mixture of numerical precisions. Traditionally, scientific computing has relied on double-precision floating point operations, however, interest in reduced numerical precision has grown in recent years due to breakthroughs in artificial intelligence and machine learning. In this case, researchers were able to implement high-speed single- and half-precision operations to gain additional performance.

The record-setting run was carried out using a representative dataset on 4,000 of Summit's GPU-accelerated nodes.

WEEK MONTH ALL TIME

Genomics Code Exceeds Exaops on Summit Supercomputer

Introduce Your Daughter to AI Event Sees OLCF Participation

International Teams Optimize Codes at Australia's First OLCF GPU Hackathon

US Air Force and ORNL Partner in High Performance Computing and Weather Modeling System

Summit by the Numbers

ESnet

Science at Scale: Climate

nature International weekly journal of science

Search Go Advanced search

Home | News & Comment | Research | Careers & Jobs | Current Issue | Archive | Audio & Video | For Authors

Archive > Volume 548 > Issue 7668 > News > Article

NATURE | NEWS

E-alert RSS Facebook Twitter

How machine learning could help to improve climate forecasts

Mixing artificial intelligence with climate science helps researchers to identify previously unknown atmospheric processes and rank climate models.

Nicola Jones

23 August 2017

PDF Rights & Permissions

Greg Kendall-Bell

Many of the latest climate models seek to increase the detail in simulations of cloud structure.

As Earth-observing satellites become more plentiful and climate models more powerful, researchers who study global warming are facing a deluge of data. Some are now turning to the latest trend in artificial intelligence (AI) to help trawl through all the information, in the hope of discovering new climate patterns and improving forecasts.

nature briefing

What matters in science — and why — free in your inbox every weekday.

Sign up

Listen

Nature Podcast

Our award-winning show features highlights from the week's edition of *Nature*, interviews with the people behind the science, and in-depth commentary and analysis from journalists around

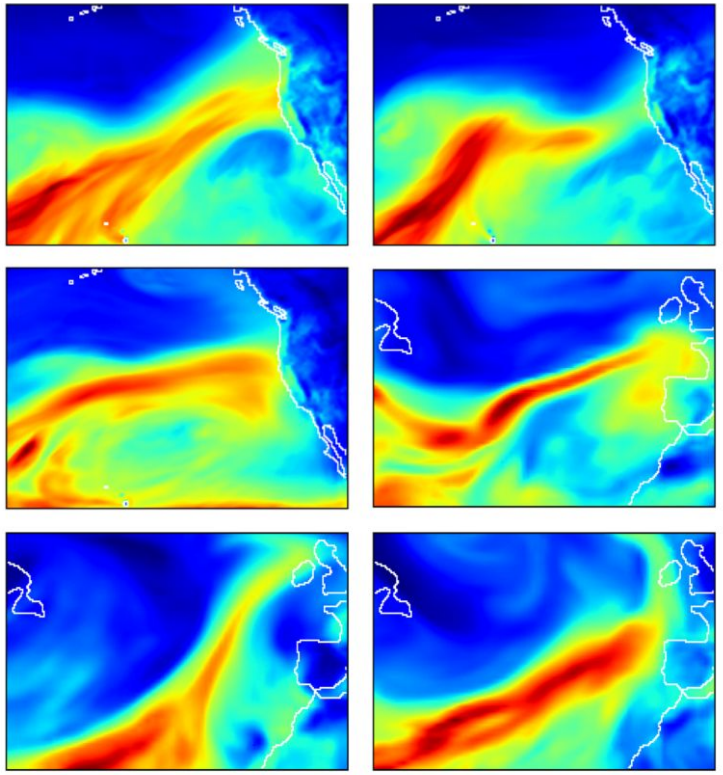

Figure 3: Sample images of atmospheric rivers correctly classified (true positive) by our deep CNN model. Figure shows total column water vapor (color map) and land sea boundary (solid line).

ESnet

They Can Use All The Data

- Groups like these need large data sets
- Much of the data in their field is behind legacy portals
  - Significant human effort to retrieve what they need
  - Legacy systems perform poorly, especially at scale
- Legacy data portals are a product of their time
  - Remember: these were designed to serve small data to small systems
  - We now live in the future from the perspective of those designs
  - Current systems far exceed the capabilities available 15 years ago
  - From the perspective of today's systems, legacy portals are products of a bygone past
- It is now perfectly reasonable for a scientist to want all the data
  - Machine learning + HPC
  - But this only works if the scientists can get to the data at scale

ESnet

We (not the scientists) Have To Do This

- The scientific community cannot do this for themselves
- Individual researchers do not control the resources
  - Computing centers
  - Data repositories
  - Science networks
  - Our community owns these – we have to do the work
- Integration, performance engineering, interoperability
- Science Engagement to teach scientists how to use the better platforms
- This is the path forward

ESnet

Networks Cannot Do This Alone

- We need a whole-community effort
  - Networks
  - HPC facilities
  - Data repositories / Data portals
  - Experimental facilities
  - Science collaborations
  - Science programs

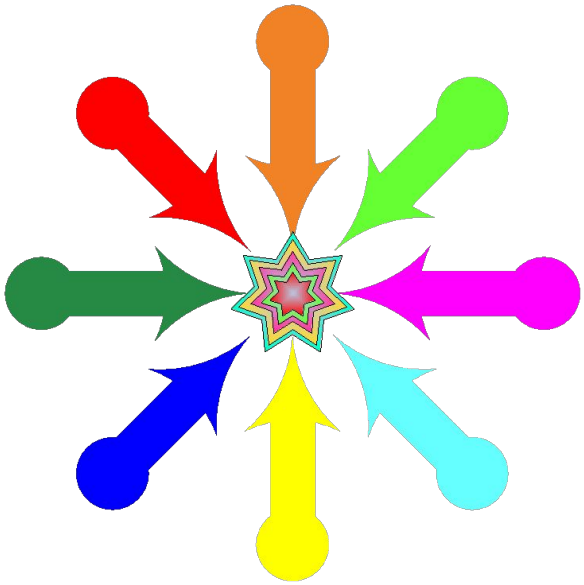

- Networks can help, and must be part of the conversation
  - Heavy lifting is now at the network edge, in collaboration with the network core
  - Need to help them get the architecture right – we know how to do this

ESnet

Vision – Interoperable Computing And Data

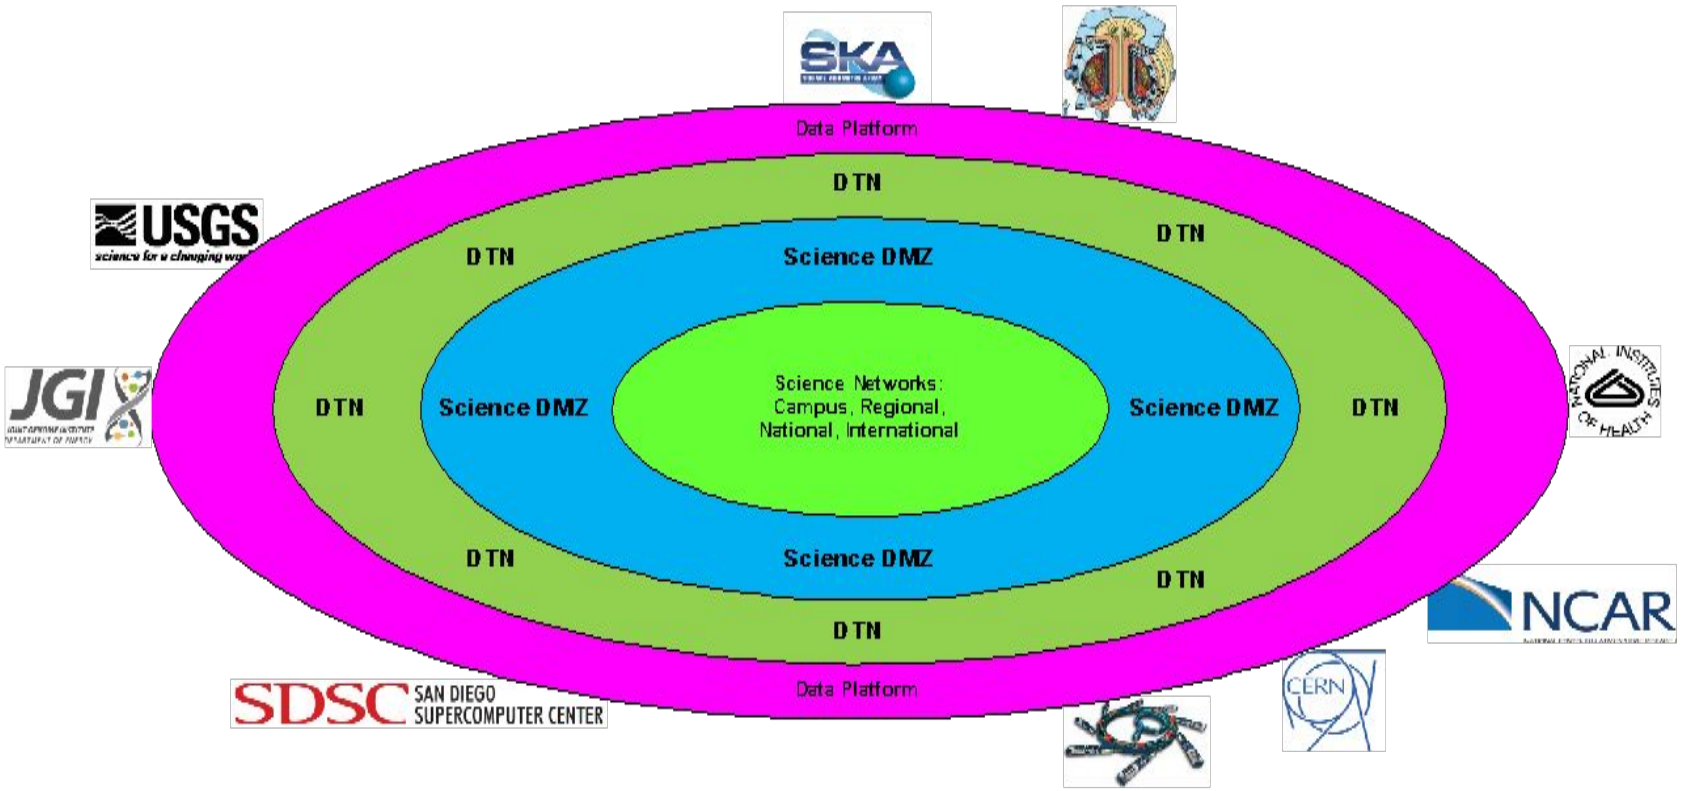

ESnet

Cyberinfrastructure Is The Ladder Up

- An integrated, interoperable cyberinfrastructure will allow scientists to make effective use of data, computing, and networks
- This is how we will achieve the advances we need in medicine, energy, climate science, and many other fields
- Large-scale data can only be effectively used if the tools work well together - otherwise the effort is too great

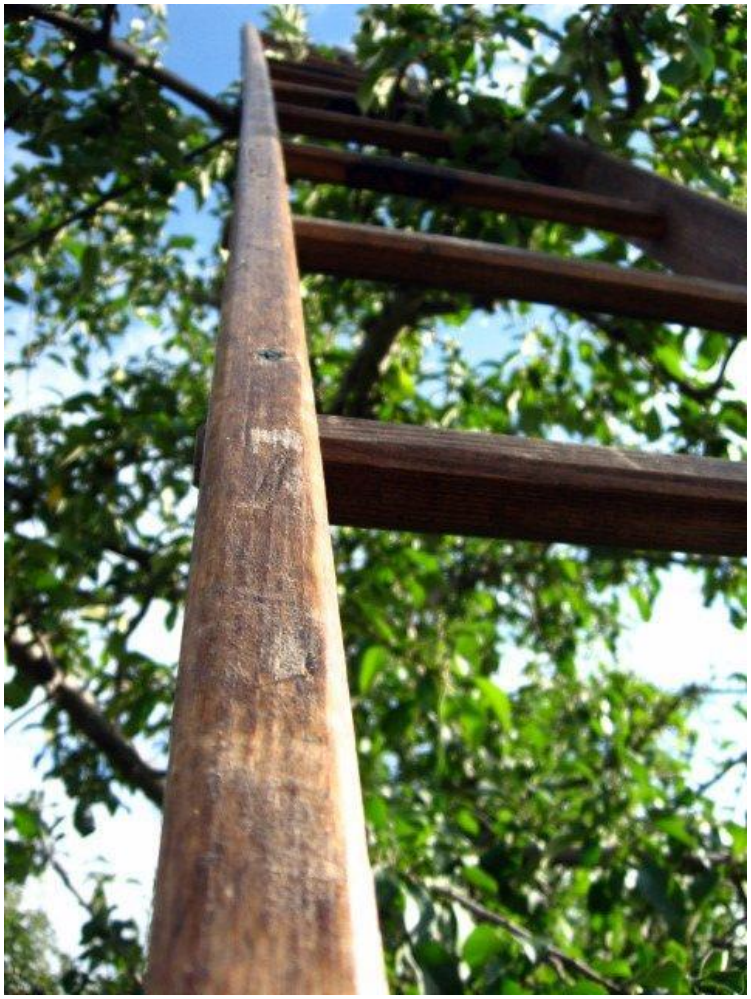

ESnet

The Path Forward

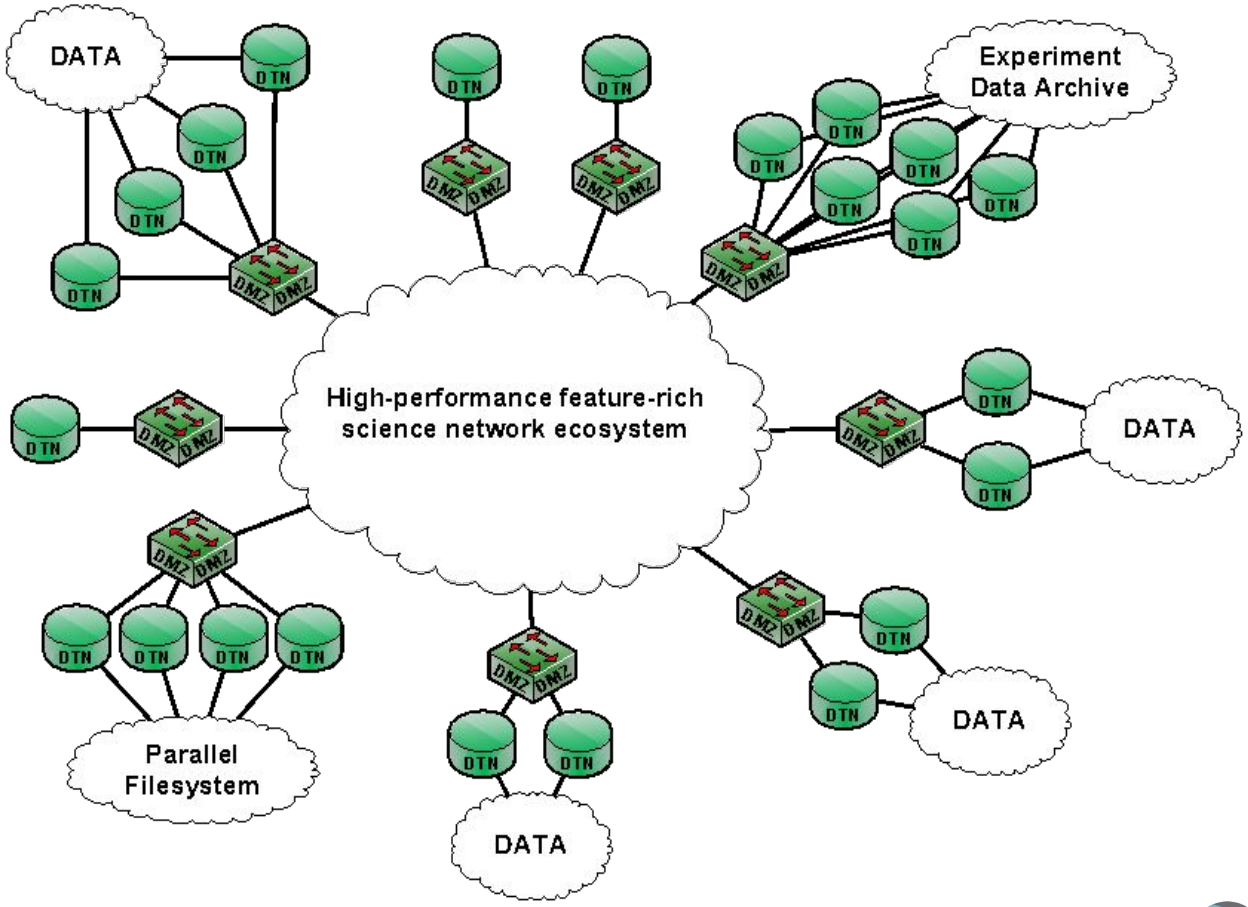

The Path Forward

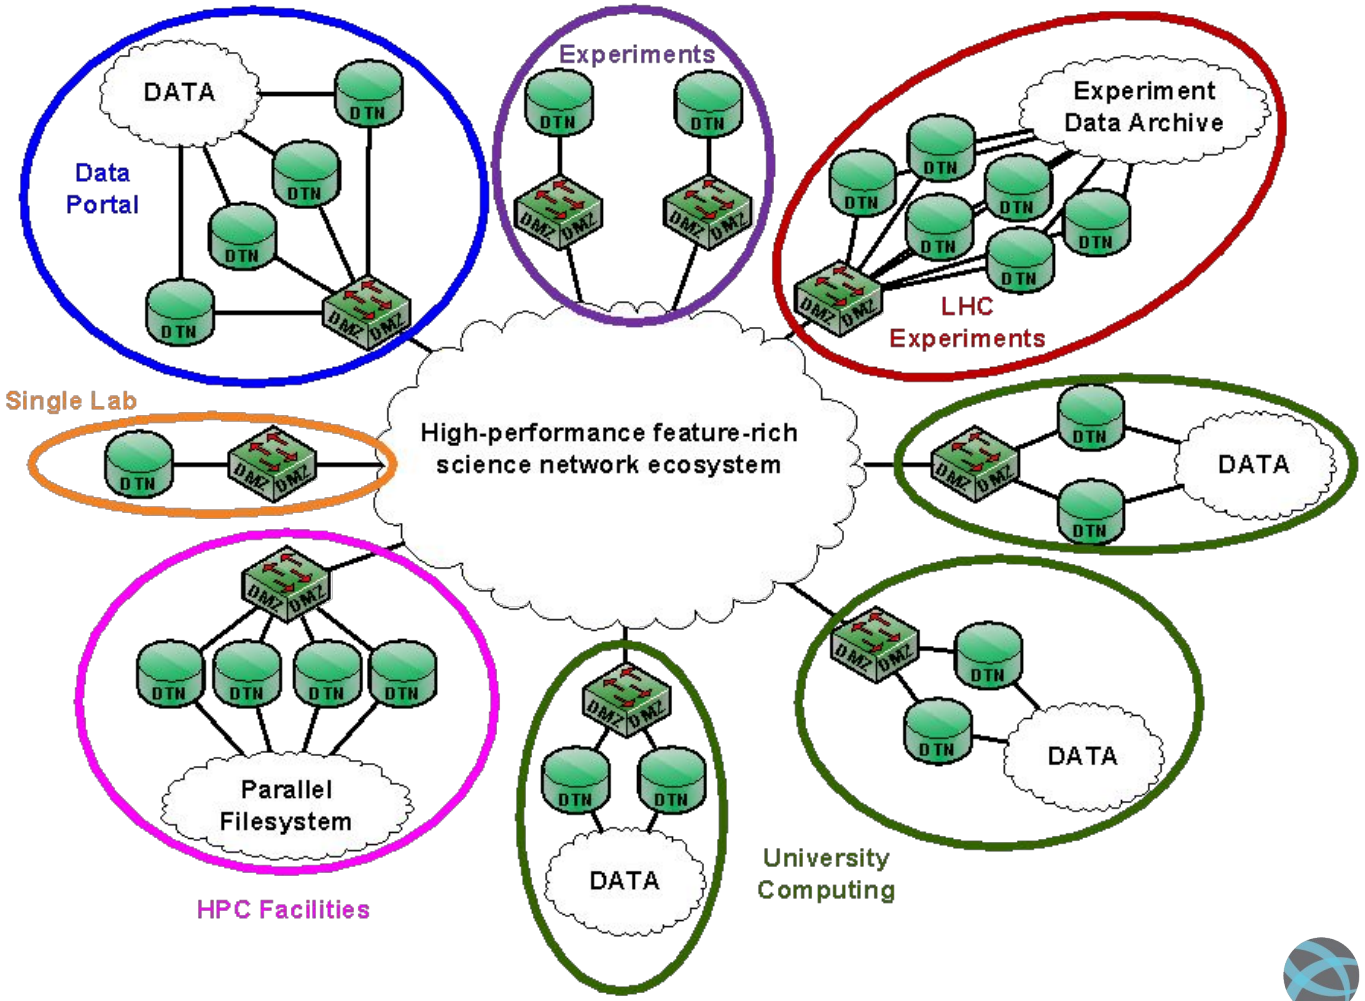

ESnet

Standing On The Shoulders Of Giants

- Large-scale data sets are the giants of today
- We have all the components we need to give all scientists access to all the data in their fields
- This is not a design problem
  - We have the designs, the technologies, the models
  - We know it works: we have examples
- This is an integration and deployment problem
  - We know what we need to do
  - Let's get to it!

Isaac Newton

ESnet

Image source: https://commons.wikimedia.org/wiki/File:Sir_Isaac_Newton_by_Sir_Godfrey_Kneller,_Bt.jpg
Godfrey Kneller [Public domain], via Wikimedia Commons

ESnet

ENERGY SCIENCES NETWORK

Thanks!

Eli Dart dart@es.net

Energy Sciences Network (ESnet)

Lawrence Berkeley National Laboratory

engage@es.net

http://my.es.net/

http://www.es.net/

http://fasterdata.es.net/

U.S. DEPARTMENT OF ENERGY

Office of Science

BERKELEY LAB

Extra slides – data download from portal

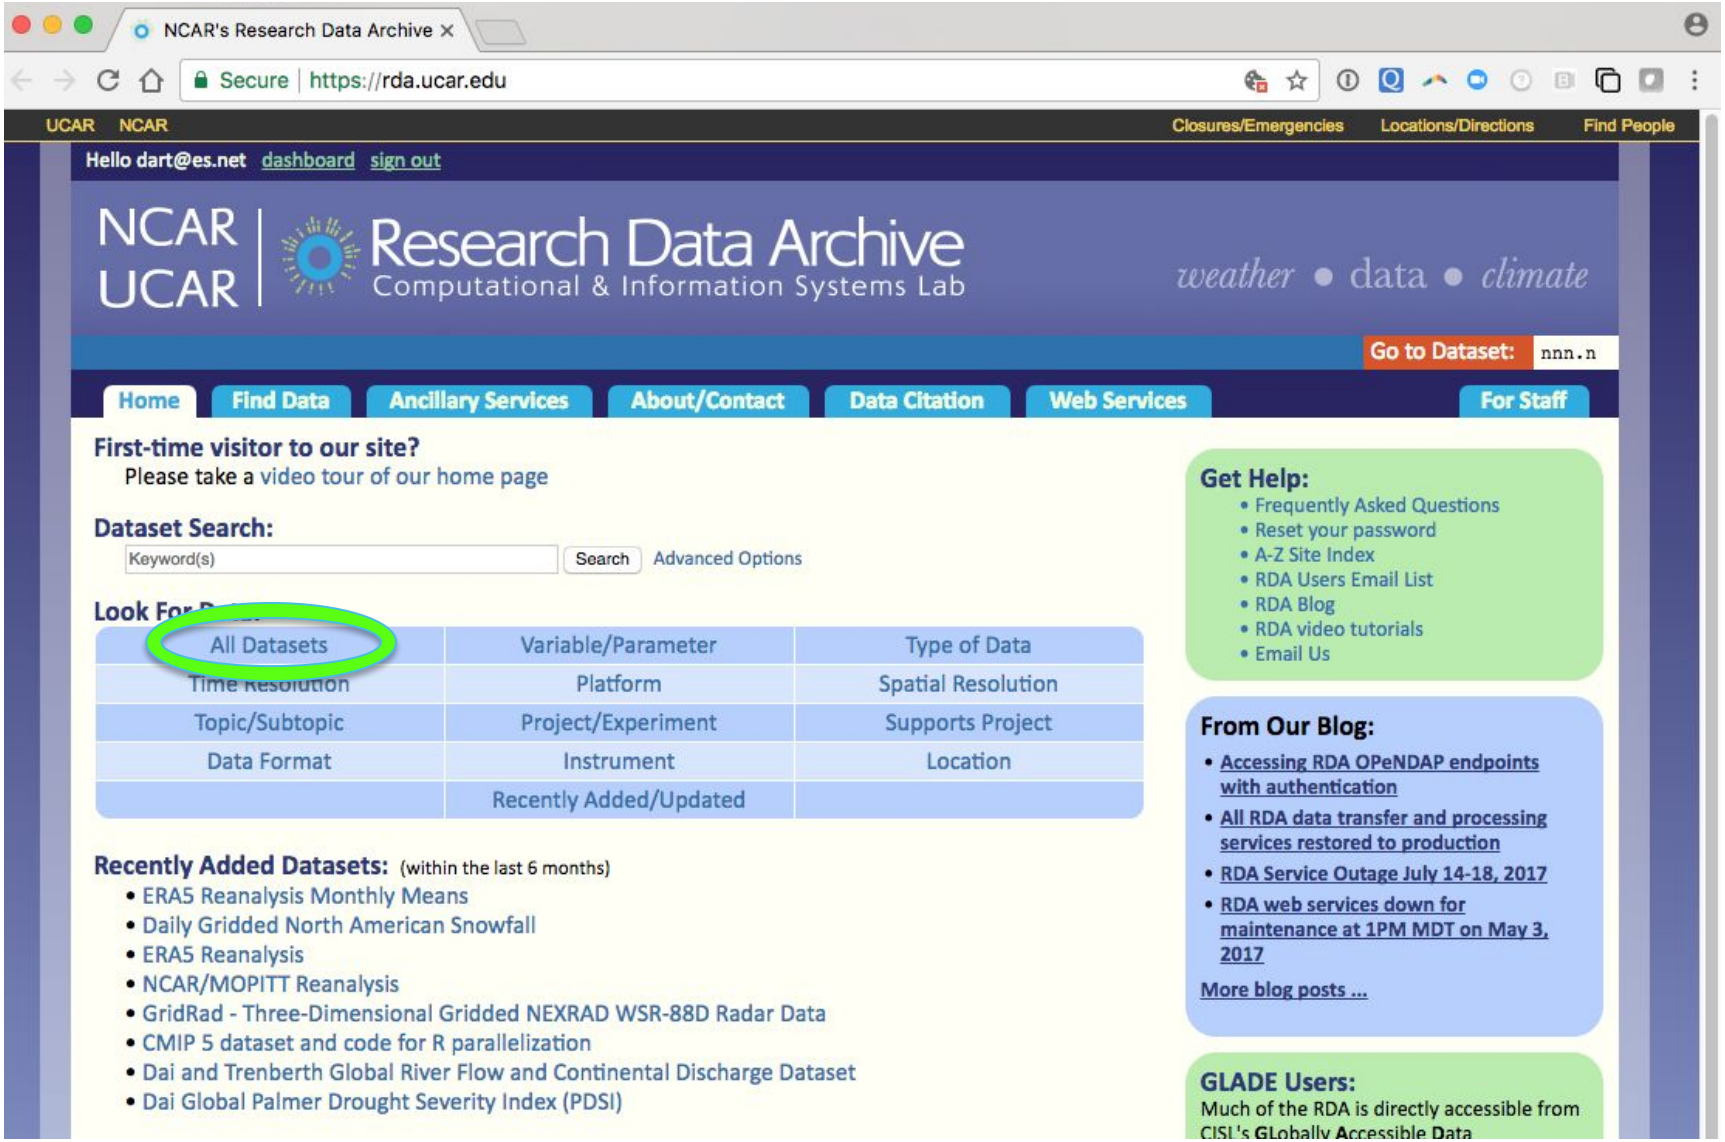
NCAR's Research Data Archive
Secure | https://rda.ucar.edu
UCAR NCAR
Closures/Emergencies
Locations/Directions
Find People
Hello dart@es.net dashboard sign out
NCAR UCAR
Research Data Archive
Computational & Information Systems Lab
weather • data • climate
Go to Dataset: nnn.n
Home
Find Data
Ancillary Services
About/Contact
Data Citation
Web Services
For Staff
First-time visitor to our site?
Please take a video tour of our home page
Dataset Search:
Keyword(s)
Search
Advanced Options
Look For
All Datasets
Variable/Parameter
Type of Data
Time Resolution
Platform
Spatial Resolution
Topic/Subtopic
Project/Experiment
Supports Project
Data Format
Instrument
Location
Recently Added/Updated
Recently Added Datasets: (within the last 6 months)
• ERA5 Reanalysis Monthly Means
• Daily Gridded North American Snowfall
• ERA5 Reanalysis
• NCAR/MOPITT Reanalysis
• GridRad - Three-Dimensional Gridded NEXRAD WSR-88D Radar Data
• CMIP 5 dataset and code for R parallelization
• Dai and Trenberth Global River Flow and Continental Discharge Dataset
• Dai Global Palmer Drought Severity Index (PDSI)
Get Help:
• Frequently Asked Questions
• Reset your password
• A-Z Site Index
• RDA Users Email List
• RDA Blog
• RDA video tutorials
• Email Us
From Our Blog:
• Accessing RDA OPeNDAP endpoints with authentication
• All RDA data transfer and processing services restored to production
• RDA Service Outage July 14-18, 2017
• RDA web services down for maintenance at 1PM MDT on May 3, 2017
More blog posts ...
GLADE Users:
Much of the RDA is directly accessible from

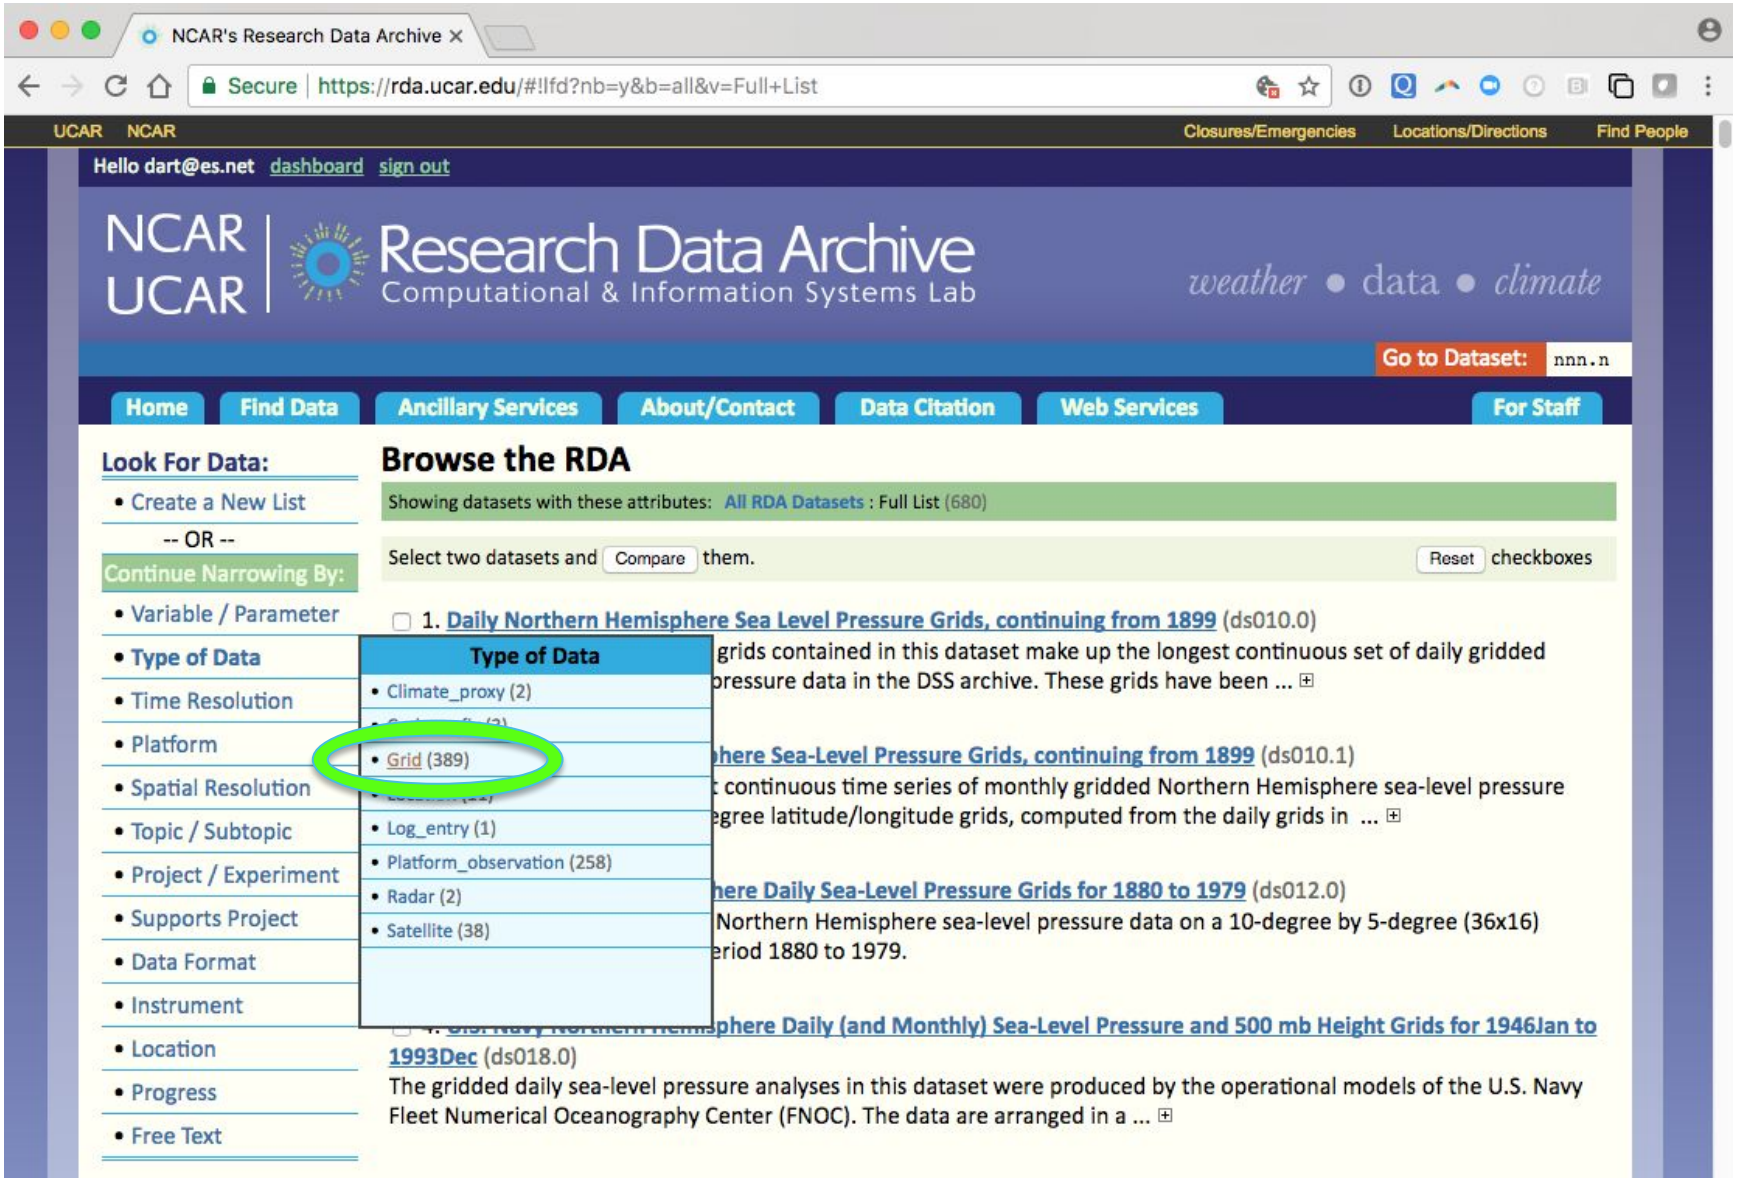
NCAR's Research Data Archive
Secure | https://rda.ucar.edu/#!lfd?nb=y&b=all&v=Full+List
UCAR NCAR
Closures/Emergencies
Locations/Directions
Find People
Hello dart@es.net dashboard sign out
NCAR UCAR
Research Data Archive
Computational & Information Systems Lab
weather • data • climate
Go to Dataset: nnn.n
Home
Find Data
Ancillary Services
About/Contact
Data Citation
Web Services
For Staff
Look For Data:
• Create a New List
-- OR --
Continue Narrowing By:
• Variable / Parameter
• Type of Data
• Time Resolution
• Platform
• Spatial Resolution
• Topic / Subtopic
• Project / Experiment
• Supports Project
• Data Format
• Instrument
• Location
• Progress
• Free Text
Browse the RDA
Showing datasets with these attributes: All RDA Datasets : Full List (680)
Select two datasets and Compare them.
Reset checkboxes
1. Daily Northern Hemisphere Sea Level Pressure Grids, continuing from 1899 (ds010.0)
grids contained in this dataset make up the longest continuous set of daily gridded
pressure data in the DSS archive. These grids have been ...
Type of Data
• Climate_proxy (2)
• Grid (389)
• Log_entry (1)
• Platform_observation (258)
• Radar (2)
• Satellite (38)
here Sea-Level Pressure Grids, continuing from 1899 (ds010.1)
continuous time series of monthly gridded Northern Hemisphere sea-level pressure
egree latitude/longitude grids, computed from the daily grids in ...
here Daily Sea-Level Pressure Grids for 1880 to 1979 (ds012.0)
Northern Hemisphere sea-level pressure data on a 10-degree by 5-degree (36x16)
eriod 1880 to 1979.
sphere Daily (and Monthly) Sea-Level Pressure and 500 mb Height Grids for 1946Jan to 1993Dec (ds018.0)
The gridded daily sea-level pressure analyses in this dataset were produced by the operational models of the U.S. Navy Fleet Numerical Oceanography Center (FNOC). The data are arranged in a ...

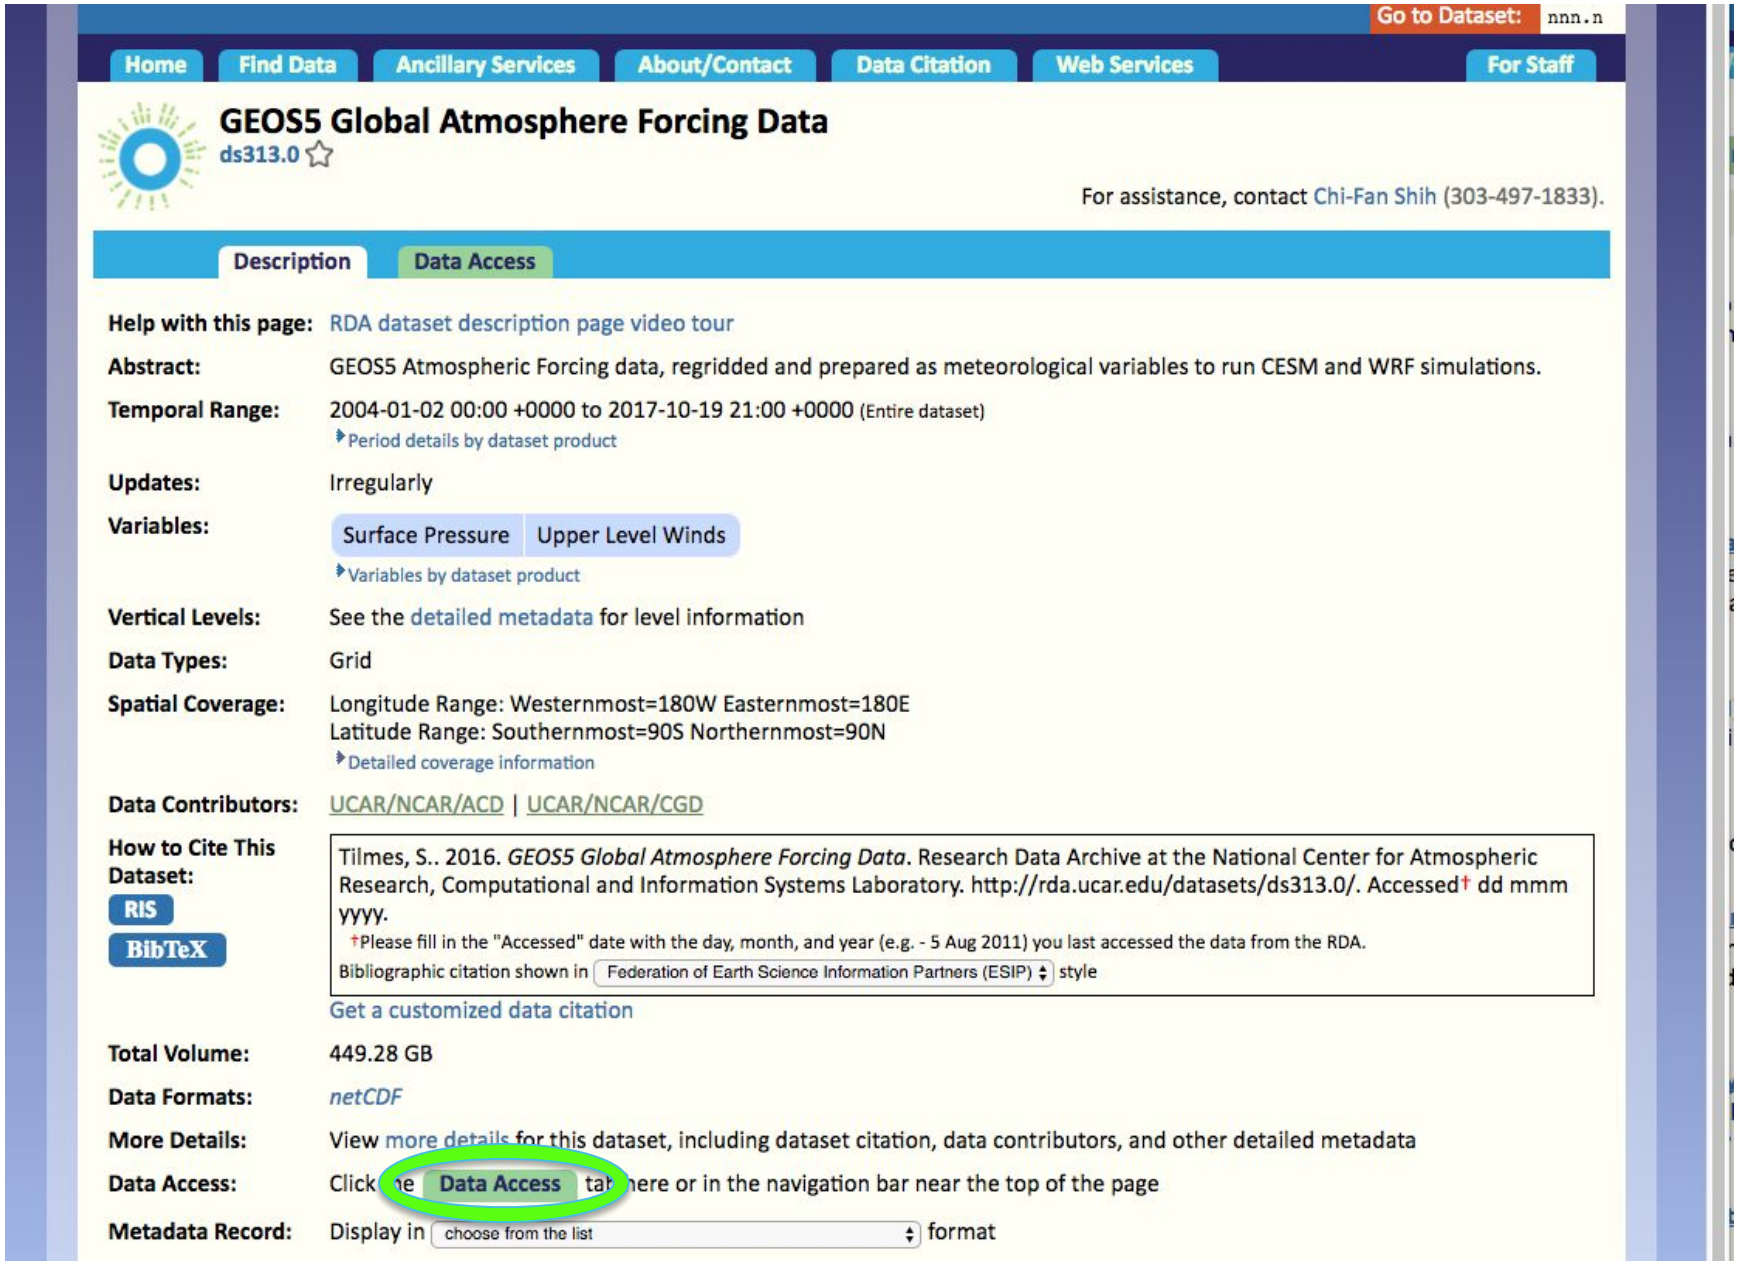
Go to Dataset: nnn.n
Home
Find Data
Ancillary Services
About/Contact
Data Citation
Web Services
For Staff
GEOS5 Global Atmosphere Forcing Data
ds313.0
For assistance, contact Chi-Fan Shih (303-497-1833).
Description
Data Access
Help with this page: RDA dataset description page video tour
Abstract: GEOS5 Atmospheric Forcing data, regridded and prepared as meteorological variables to run CESM and WRF simulations.
Temporal Range: 2004-01-02 00:00 +0000 to 2017-10-19 21:00 +0000 (Entire dataset)
Period details by dataset product
Updates: Irregularly
Variables: Surface Pressure Upper Level Winds
Variables by dataset product
Vertical Levels: See the detailed metadata for level information
Data Types: Grid
Spatial Coverage: Longitude Range: Westernmost=180W Easternmost=180E
Latitude Range: Southernmost=90S Northernmost=90N
Detailed coverage information
Data Contributors: UCAR/NCAR/ACD | UCAR/NCAR/CGD
How to Cite This Dataset:
RIS
BibTeX
Tilmes, S.. 2016. GEOS5 Global Atmosphere Forcing Data. Research Data Archive at the National Center for Atmospheric Research, Computational and Information Systems Laboratory. http://rda.ucar.edu/datasets/ds313.0/. Accessed† dd mmm yyyy.
†Please fill in the "Accessed" date with the day, month, and year (e.g. - 5 Aug 2011) you last accessed the data from the RDA.
Bibliographic citation shown in Federation of Earth Science Information Partners (ESIP) style
Get a customized data citation
Total Volume: 449.28 GB
Data Formats: netCDF
More Details: View more details for this dataset, including dataset citation, data contributors, and other detailed metadata
Data Access: Click Data Access tab here or in the navigation bar near the top of the page
Metadata Record: Display in choose from the list format

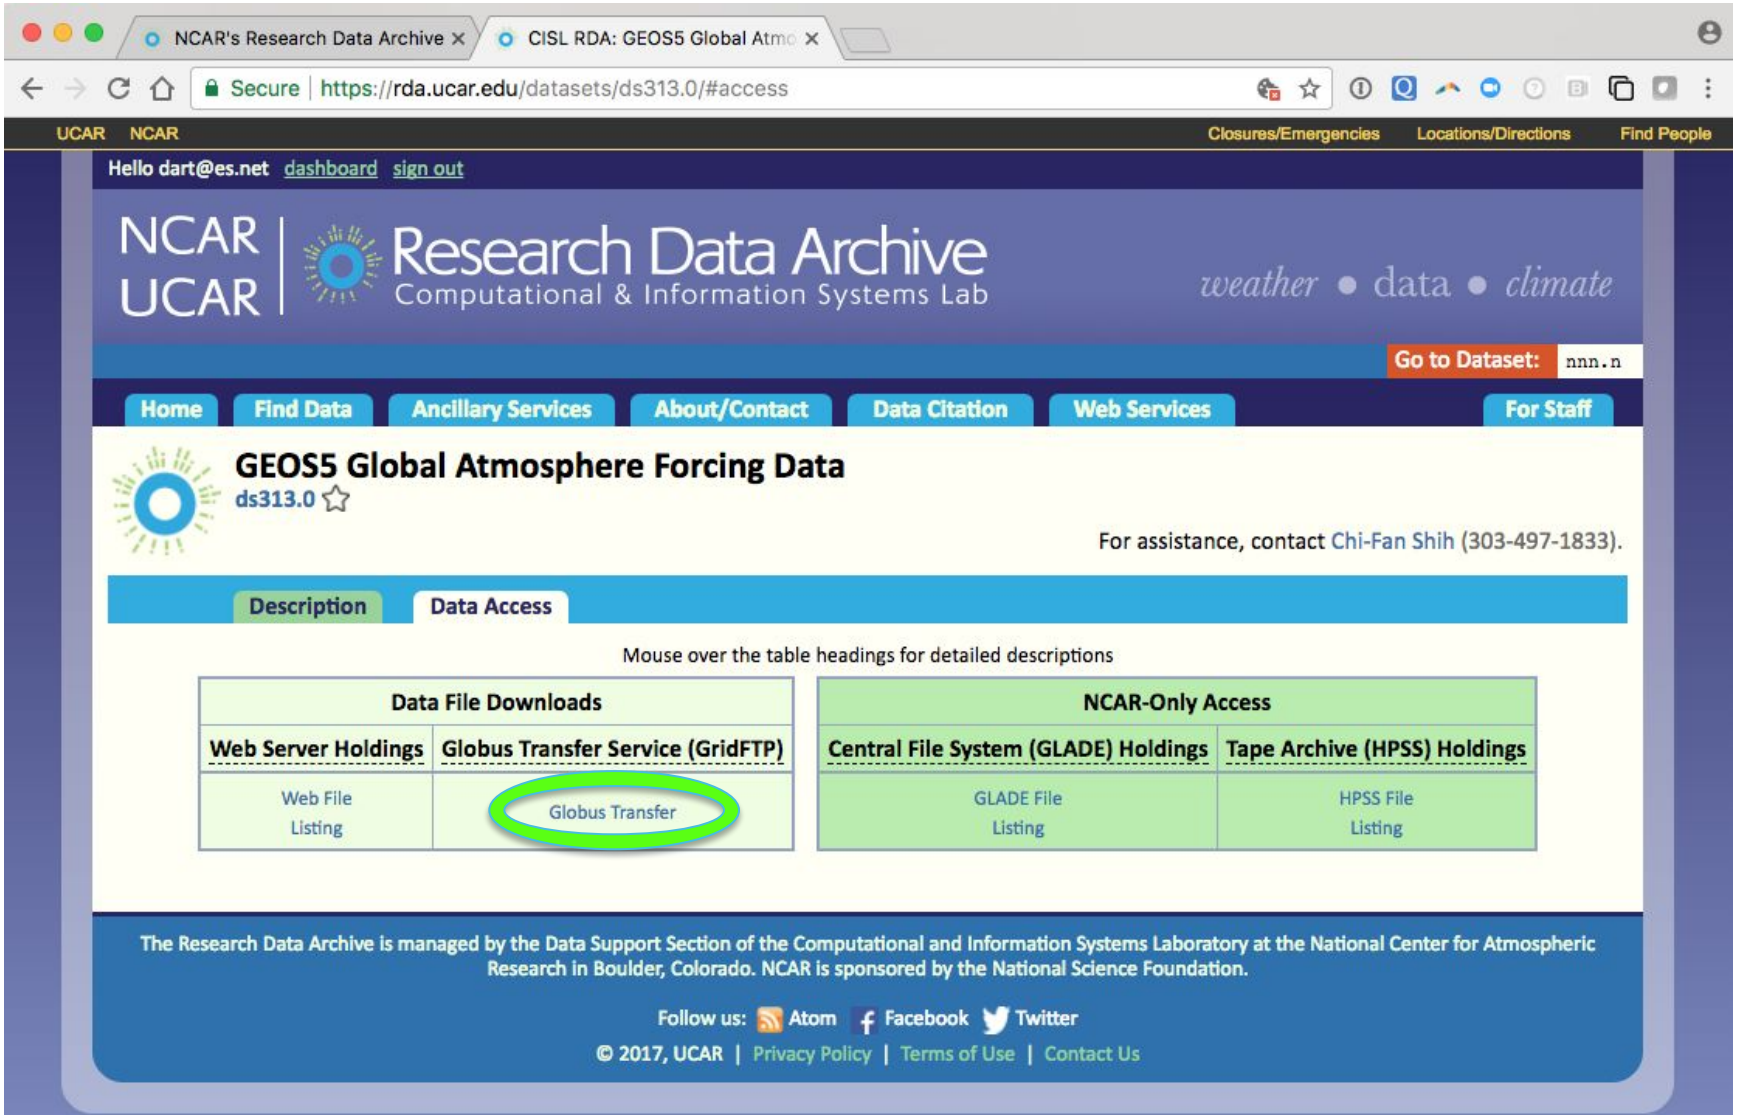
NCAR's Research Data Archive
CISL RDA: GEOS5 Global Atmo
Secure | https://rda.ucar.edu/datasets/ds313.0/#access
UCAR NCAR
Closures/Emergencies
Locations/Directions
Find People
Hello dart@es.net dashboard sign out
NCAR UCAR
Research Data Archive
Computational & Information Systems Lab
weather • data • climate
Go to Dataset: nnn.n
Home
Find Data
Ancillary Services
About/Contact
Data Citation
Web Services
For Staff
GEOS5 Global Atmosphere Forcing Data
ds313.0
For assistance, contact Chi-Fan Shih (303-497-1833).
Description
Data Access
Mouse over the table headings for detailed descriptions
Data File Downloads
Web Server Holdings
Globus Transfer Service (GridFTP)
Web File Listing
Globus Transfer
NCAR-Only Access
Central File System (GLADE) Holdings
Tape Archive (HPSS) Holdings
GLADE File Listing
HPSS File Listing
The Research Data Archive is managed by the Data Support Section of the Computational and Information Systems Laboratory at the National Center for Atmospheric Research in Boulder, Colorado. NCAR is sponsored by the National Science Foundation.
Follow us: Atom Facebook Twitter
© 2017, UCAR | Privacy Policy | Terms of Use | Contact Us

Portal creates a Globus transfer job for us

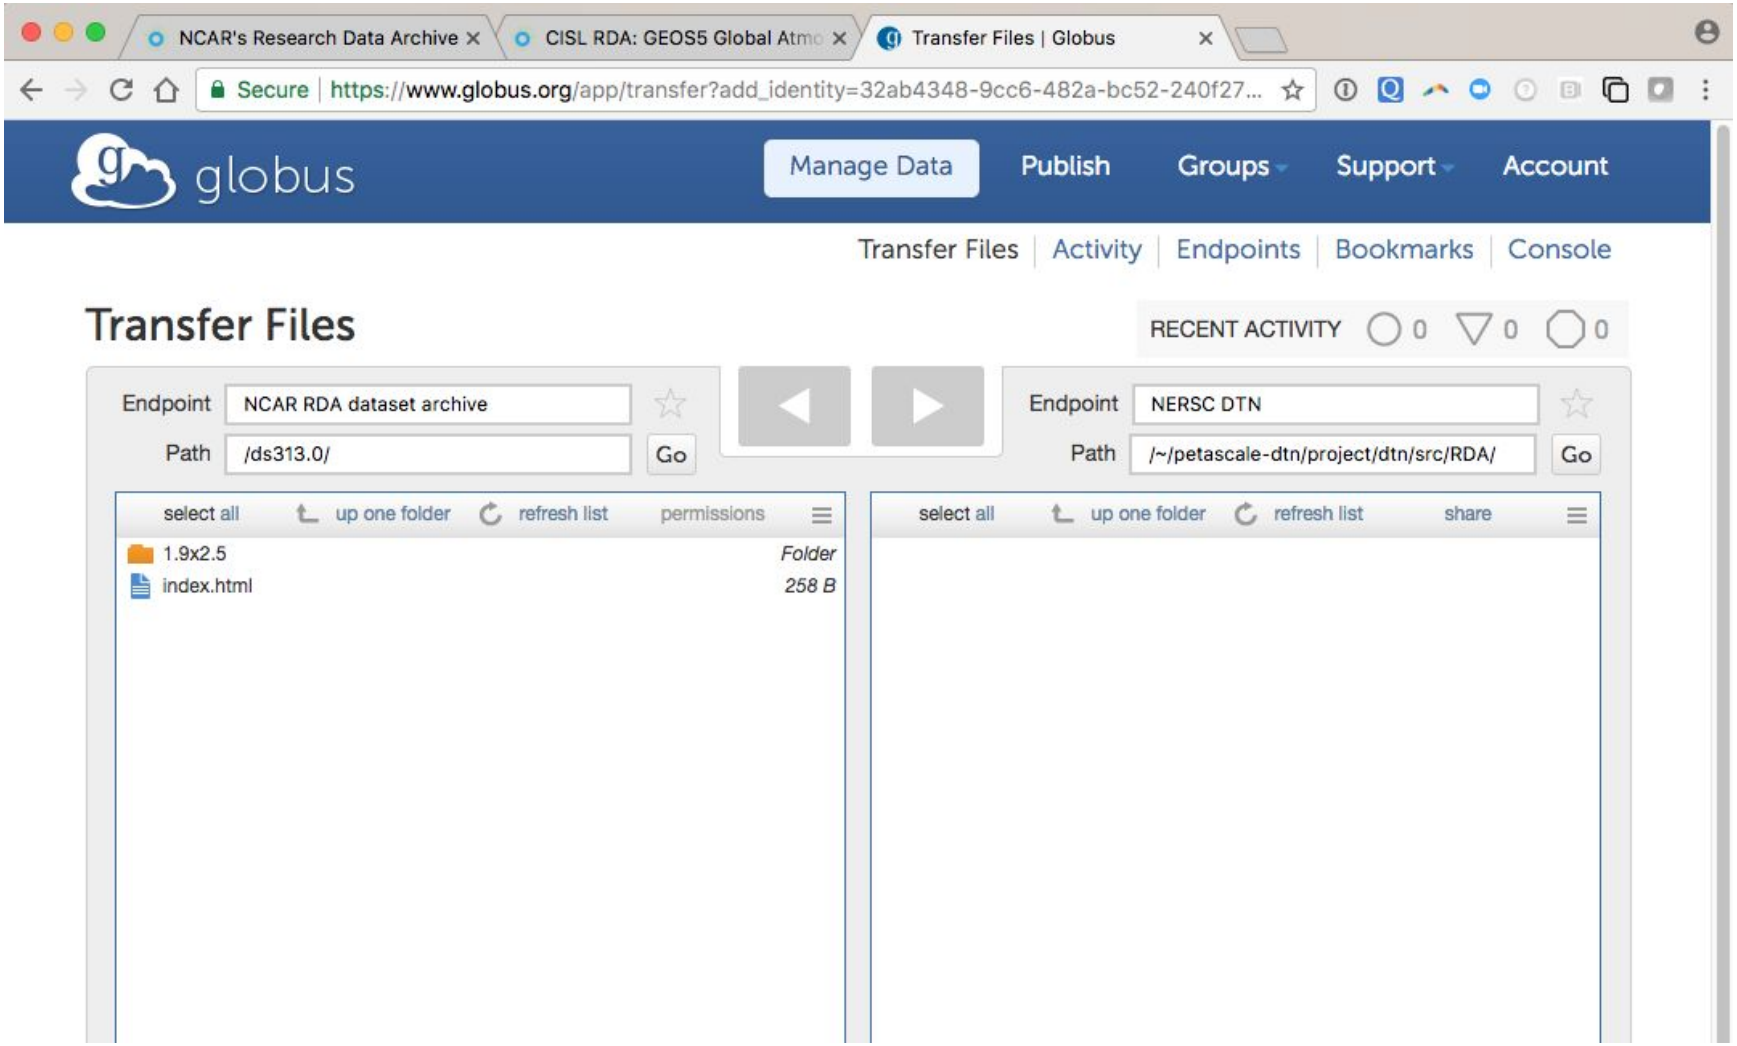

Submit the transfer job, go about our business

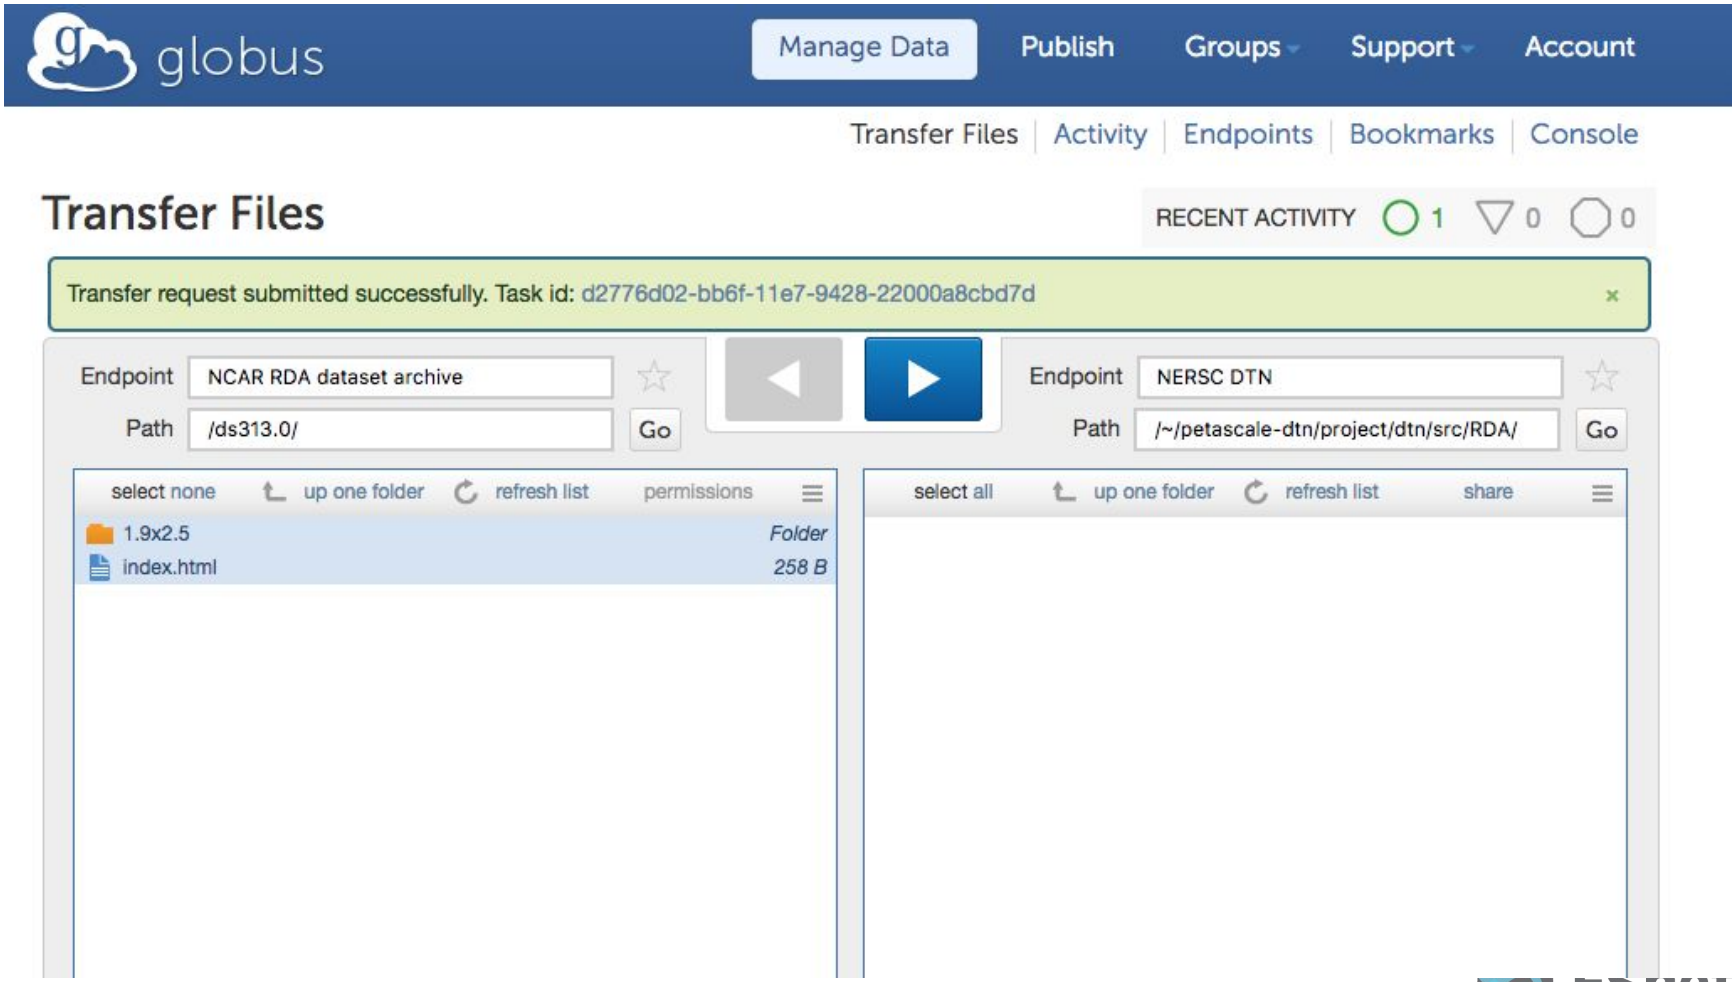

Data Transfer from RDA Portal – Results

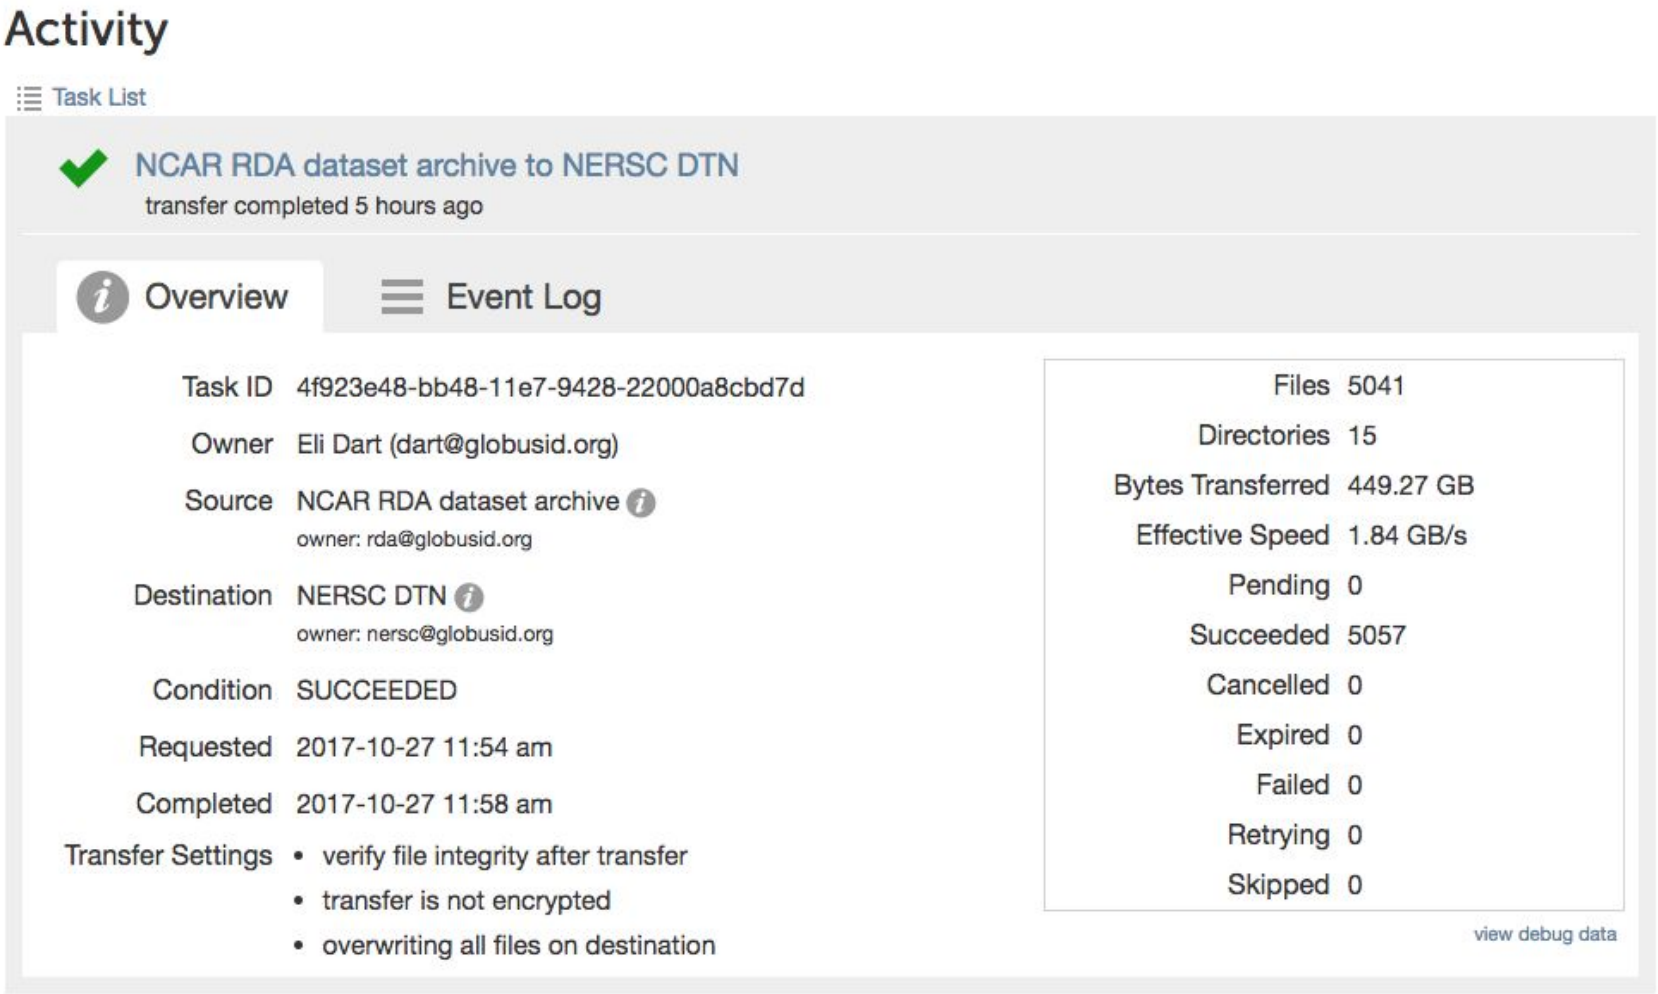

Superfacility: Integrated network of experimental and computational facilities and expertise

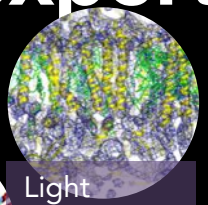

Computing and Data Facilities

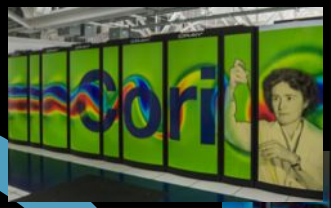

ESnet

Science Network

Experimental Facilities

A single interconnected "facility" where data is acquired, stored, analyzed and served

Expertise

Methods, models, analytics, and software

More Users

More Productive Science

Advanced Light Source Demonstration

ALS ADVANCED LIGHT SOURCE

ESnet

GISAXS

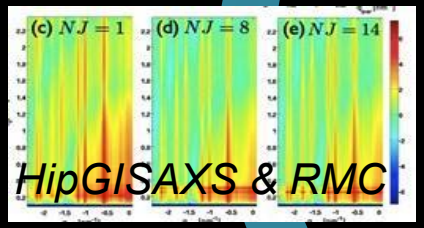

CAMERA Applied Math

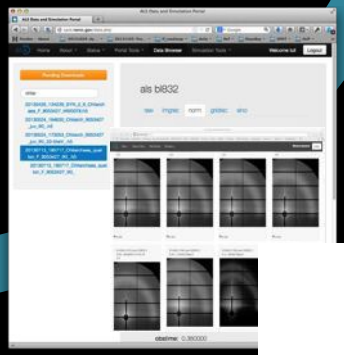

Liu et al, “Fast printing and in situ morphology …”. Adv Mater. 2015